



---

Serverless Messaging  
with Pub/Sub

# Agenda

---

Processing Streaming Data

Cloud Pub/Sub

Cloud Dataflow Streaming  
Features

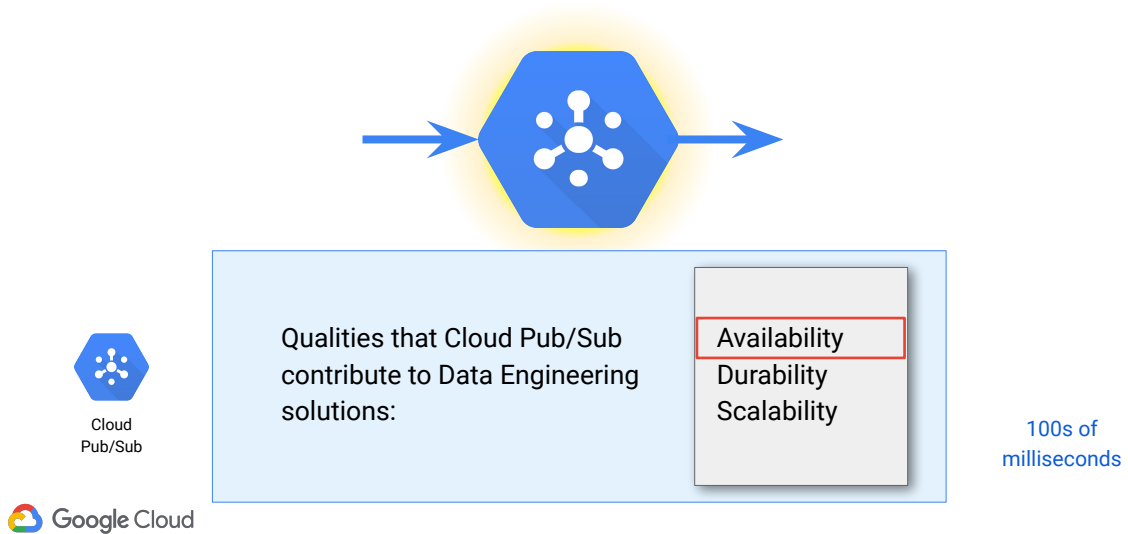
BigQuery and Bigtable Streaming  
Features

Advanced BigQuery  
Functionality



Now, that we have a good understanding of the process of streaming data, let's dive into Cloud Pub/Sub to see how it works. As we begin this topic, I would ask you to keep your mind open to new ways of doing things. Cloud Pub/Sub does streaming differently than probably anything you have used in the past. It may be a different model than what you have seen before.

# Cloud Pub/Sub



Cloud Pub/Sub provides a fully managed data distribution and delivery system. It can be used for many purposes. It is most popularly used to loosely-couple parts of a system. You can use Cloud Pub/Sub to connect applications within Google Cloud, and with applications on premise or in other clouds to create hybrid Data Engineering solutions. With Cloud Pub/Sub the applications do not need to be online and available all the time. And the parts do not need to know how to communicate to each other, but only to Cloud Pub/Sub, which can simplify system design.

First, Pub/Sub is not software; it is a service. So, like all of the other serverless services we have looked at, you don't install anything to use Pub/Sub. It is totally a service.

Cloud Pub/Sub client libraries are available in C#, GO, Java, Node.js, Python, and Ruby. These wrap REST API calls which can be made in any language.

It is highly available.

## Cloud Pub/Sub



Cloud  
Pub/Sub

Qualities that Cloud Pub/Sub  
contribute to Data Engineering  
solutions:

Availability  
Durability  
Scalability

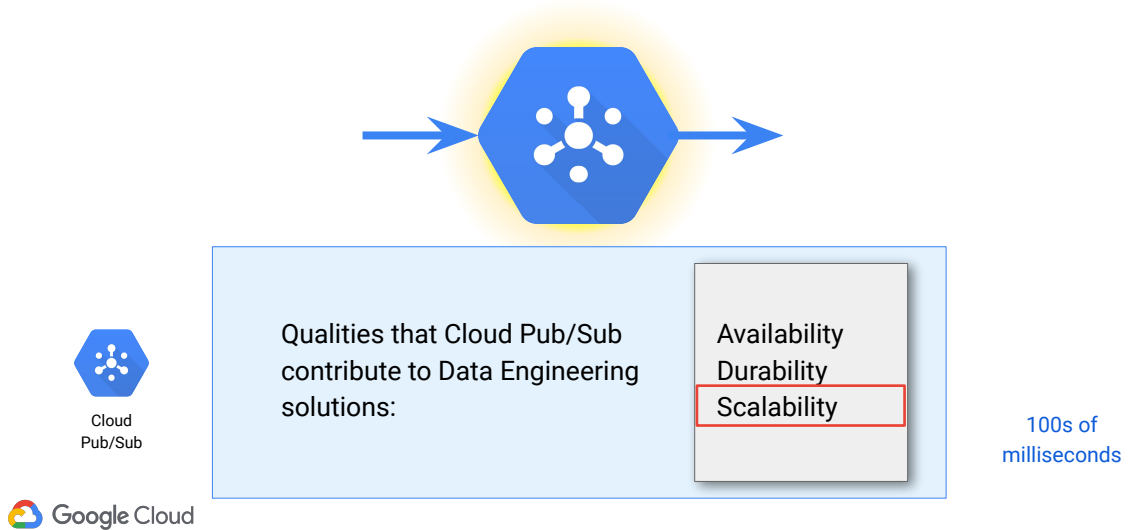
100s of  
milliseconds



It offers durability of messages. By default it will save your messages for seven days. In the event your systems are down and not able to process them.

<https://pixabay.com/illustrations/chess-black-and-white-pieces-3413429/>

# Cloud Pub/Sub



Finally, Cloud Pub/Sub is highly scalable as well. We, Google, actually process internally, about 100 million messages per second across the entire infrastructure. This was actually one of the use cases for Pub/Sub early on at Google, to be able to distribute the search engine and index around the world because we keep local copies of search around the world, as you might imagine, in order to be able to serve up results with minimal latency.

If you think about how you would architect that, we are crawling the entire world wide web, so we need to send the entire world wide web around the world. Either that or we would need to have multiple crawlers all over the world, but then you have data consistency problems; they would all be getting different indexes.

Therefore, what we do is use Pub/Sub to distribute. As the crawler goes out, it grabs every page from the world wide web, and we send every single page as a message on Pub/Sub and it gets picked up by all local copies of the search index so it can be indexed.

Currently, Google indexes the web anywhere from every two weeks, which is the slowest, to more than once an hour, for example on really popular news sites. So, on average, Google is indexing the web, on average, three times a day. Thus, what we are doing is sending the entire world wide web over Pub/Sub three times a day.

This demonstrates how Pub/Sub scales.

# Cloud Pub/Sub



Cloud  
Pub/Sub

Qualities that Cloud Pub/Sub  
contribute to Data Engineering  
solutions:

Availability  
Durability  
Scalability

100s of  
milliseconds




Cloud Pub/Sub is a HIPAA-compliant service, offering fine-grained access controls and end-to-end encryption. Messages are encrypted in transit and at rest. Messages are stored in multiple locations for durability and availability.

You control the qualities of your Cloud Pub/Sub solution by the number of publishers, number of subscribers, and the size and number of messages. These factors provide tradeoffs between scale, low latency, and high throughput.

<https://pixabay.com/illustrations/chess-black-and-white-pieces-3413429/>

## Example of a Cloud Pub/Sub application



How does Pub/Sub work? The model is very simple. The story of Cloud Pub/Sub is the story of two data structures, the Topic and the Subscription. Both the Topic and the Subscription are abstract concepts that exist in the Pub/sub framework independently of any workers, subscribers, etc. The Cloud Pub/Sub client that creates the Topic is called the Publisher. And the Cloud Pub/Sub client that creates the Subscription is called the Subscriber.


In this example, the Subscription is subscribed to the Topic.

<INSTRUCTOR>

To receive messages published to a topic, you must create a subscription to that topic. Only messages published to the topic after the subscription is created are available to subscriber applications. The subscription connects the topic to a subscriber application that receives and processes messages published to the topic. A topic can have multiple subscriptions, but a given subscription belongs to a single topic.

</INSTRUCTOR>

A new employees arrives causing a new hire event.



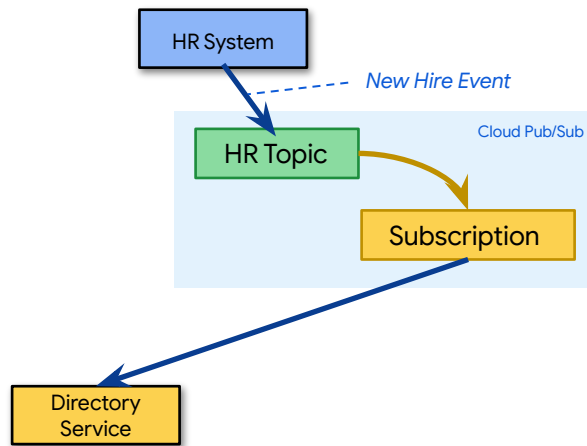
It is an enterprise message bus. So, how does it work?

Here you see, there is an HR Topic that relates to New Hire Events. For example, a new person joins your company and this notification should allow other applications that need to be notified about a new user joining to subscribe and get that message. What applications could tell you that a new person joined?

<https://pixabay.com/illustrations/avatar-clients-customers-icons-2155431/>

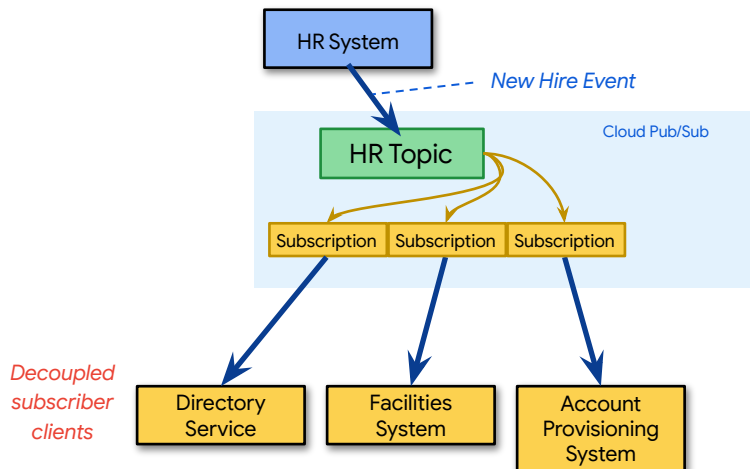


The message is sent from Topic to Subscription



One example is the Company Directory. This is a client of the Subscription also called a Subscriber. However, Cloud Pub/Sub is not limited to one Subscriber or one Subscription.

## There can be multiple Subscriptions for each Topic



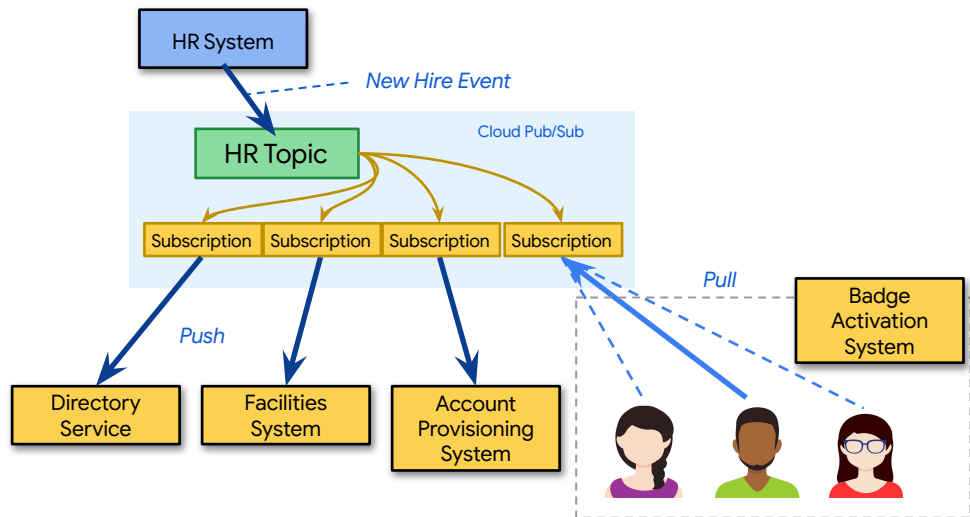
Here there are multiple Subscriptions and multiple Subscribers. Maybe the Facilities System needs to know about the new employee for badging, and the accounting provisioning system needs to know for payroll.

Each Subscriptions guarantees delivery of the message to the service.

These subscriber clients are decoupled from one another and isolated from the publisher. In fact, we will see later that the HR System could go offline after it has sent its message to the HR Topic, and the message will still be delivered to the subscribers.

These examples show one Subscription and one Subscriber. But you can actually have more Subscribers for a single Subscription.

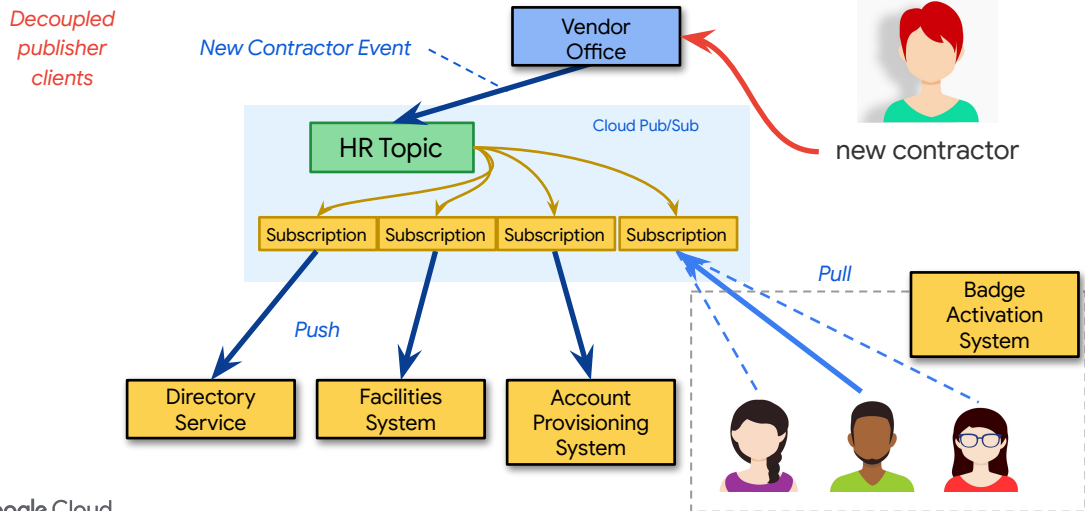
And there can be multiple subscribers per Subscription



In this example, the Badge Activation System requires a human being to activate the badge. There are multiple workers, but not all of them are available all the time. Cloud Pub/Sub makes the message available to all of them. But only one person needs to fetch the message and handle it. This is called a Pull Subscription. The other examples are Push Subscriptions.

<https://pixabay.com/illustrations/avatar-clients-customers-icons-2191918/>

## And there can be multiple publishers to the Topic



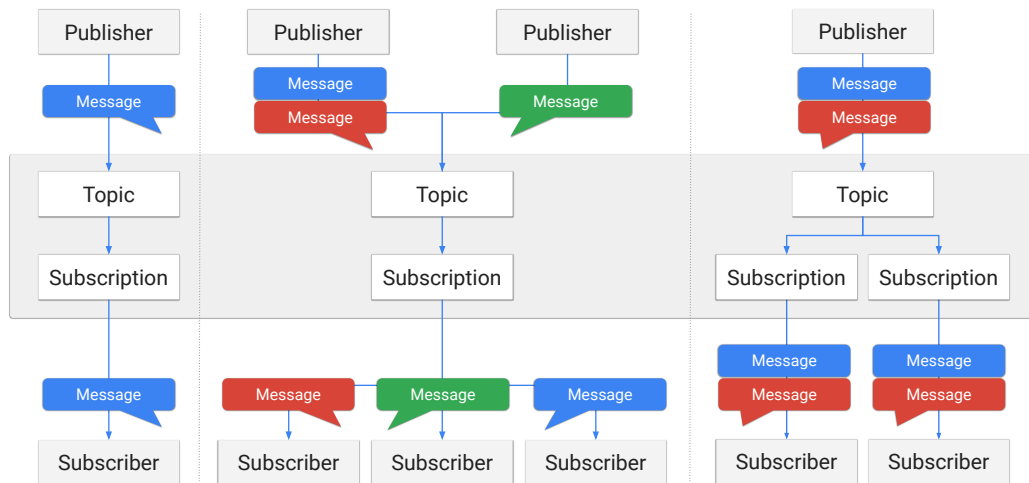
Now, a new contractor arrives. Instead of entering through the HR System, they go through the Vendor Office. The same kinds of actions need to occur for this worker. They need to be listed in the company directory. The facilities needs to assign them a desk. Account provisioning needs to set up their corporate identity and accounts. And the badge activation system needs to print and activate their contractor badge. A message can be published by the Vendor Office to the HR Topic. The Vendor Office and the HR System are entirely decoupled from one another but can make use of the same company services.

You can see from this illustration how important the Pub/Sub is. Therefore, it gets the highest priority.

You have learned generally what Cloud Pub/Sub does. Next, you will learn how it works and many of the advanced features it provides.

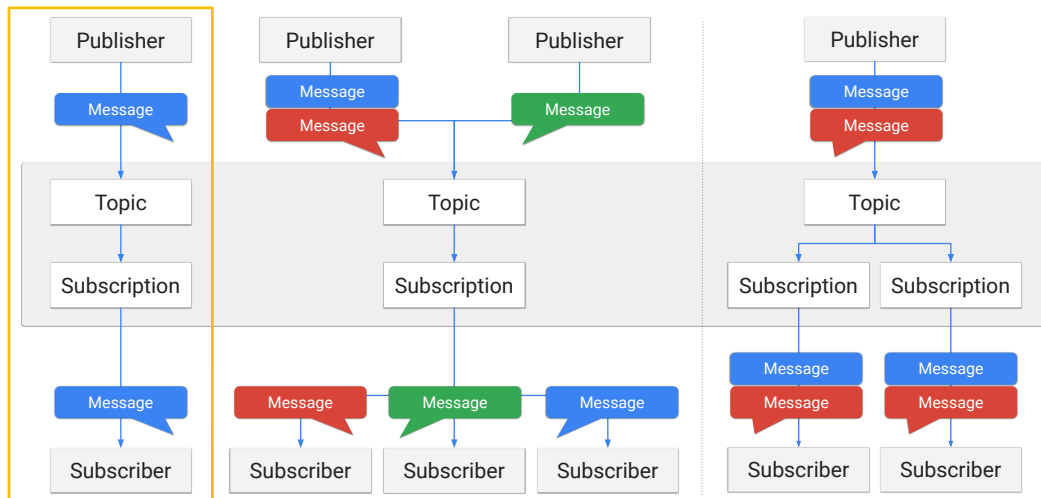
<https://pixabay.com/illustrations/avatar-clients-customers-icons-2191918/>

## Publish/Subscribe patterns



You can use this distribution of messages, the publish and subscribe patterns, to do fan in or fan out. What you see here, the different colors represent different messages.

## Publish/Subscribe patterns

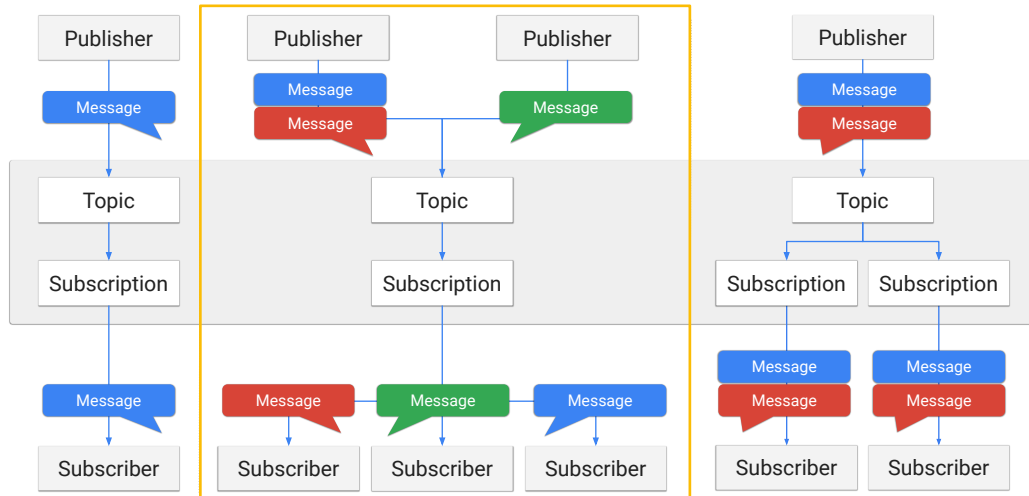


Basic pattern just a straight through, which is a queue.

(Fan in / Load Balancing) Multiple Publishers publishing to the same topic  
Also have Multiple subscribers pulling from same subscriptions, basic parallelization of processing. dataflow,  
one subscription multiple consumers and each Subscriber receives a subset of messages from the subscription.

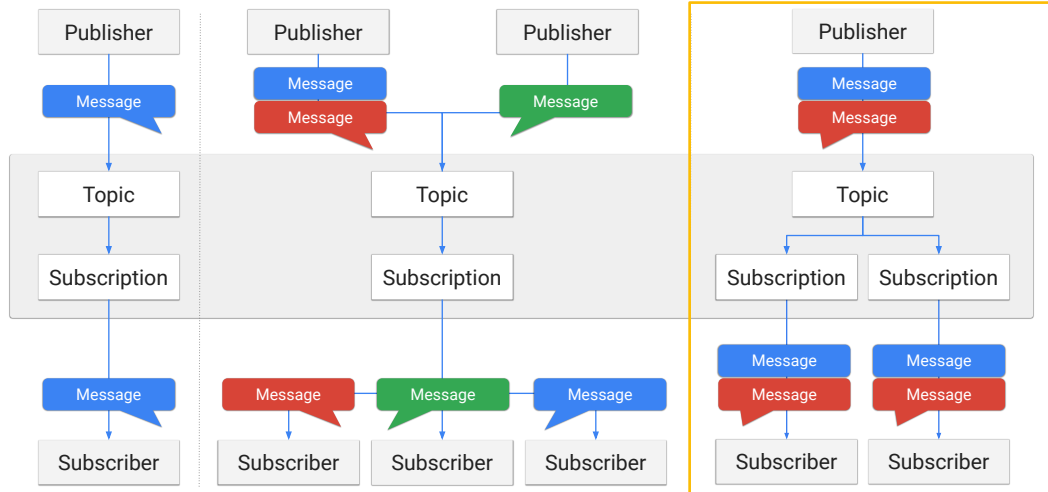
(Fan out) Multiple subscribers, where you have multiple use case for same piece of data, and all data is sent to multiple different subscribers.

## Publish/Subscribe patterns



What you see in the second example is three different messages from two different publishers being sent, but all on the same topic, and that means the subscription will get all three messages.

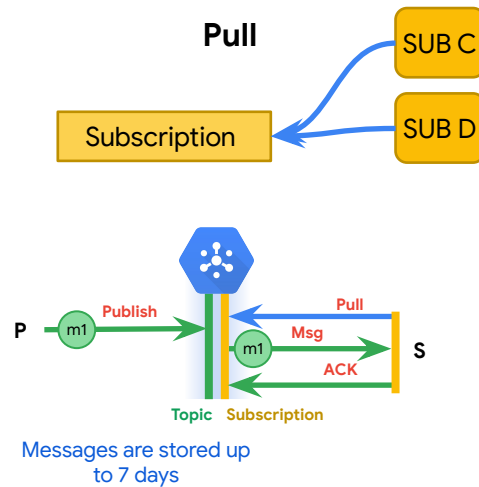
## Publish/Subscribe patterns



Over on the right, we have two subscriptions, so both are going to get the messages, both the red message and the blue message.



## Cloud Pub/Sub provides both Push and Pull delivery



Cloud Pub/Sub allows for both Push and Pull delivery. In the pull model, your clients are subscribers and will be periodically calling for messages and Pub/Sub will just be delivering the messages since the last call.

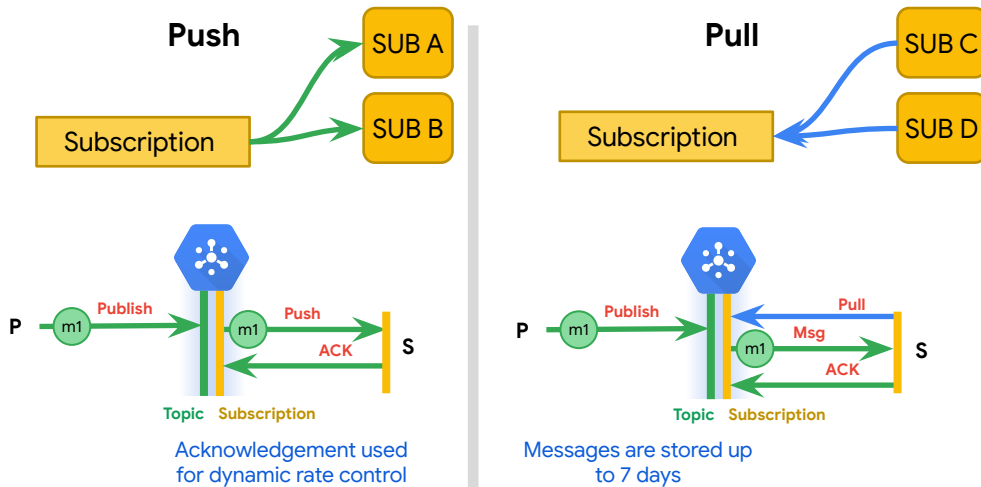
In the Pull model, you're going to have to acknowledge the message as a separate step. So, what you see here is we initially make the call to the subscribers, it pulls the messages, it gets a message back, and then, separately, it acknowledges that message.

The reason for this is because the pull queues are often used to implement some kind of queueing system for work to be done, so you don't want to acknowledge the message until you firmly have the message and have done the processing on it, otherwise you might lose the message if the system goes down.

Therefore, we generally recommend you wait to acknowledge until after you have gotten it.

In the Pull model, the messages are stored for up to seven days.

## Cloud Pub/Sub provides both Push and Pull delivery

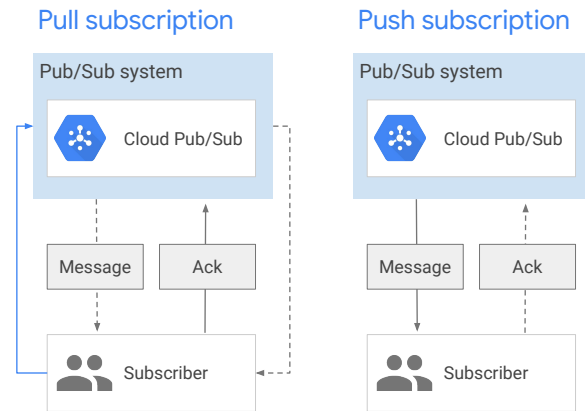


In the Push model, it actually uses an HTTP endpoint. You register an webhook as your subscription, and Pub/Sub infrastructure itself will call you with the latest messages. In the case of Push, you just respond with 'status 200 ok' for the HTTP call, and that tells Pub/Sub the message delivery was successful.

It will actually use the rate of your success responses to self limit so that it doesn't overload your worker.

## At least once delivery guarantee

- A subscriber ACKs each message for every subscription
- A message is resent if subscriber takes more than `ackDeadline` to respond
- Messages are stored for up to 7 days
- A subscriber can extend the deadline per message

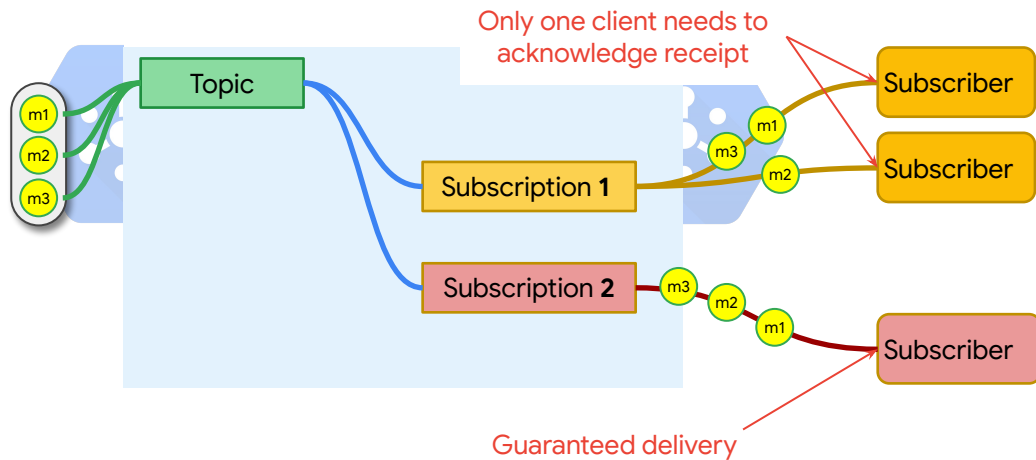


The way the acknowledgements work is to ensure every message gets delivered at least once. What happens is when you acknowledge a message, you acknowledge on a per subscription basis. So, if you have two subscriptions, you have one acknowledge and the other one doesn't, the one that acknowledged will continue to get the messages. Pub/Sub will continue to try to deliver the message up to seven days until it is acknowledged.

Now, there is a replay mechanism as well that you can rewind and go back in time and have it replay messages, but in any case, you will always be able to go back seven days.

You can also set the acknowledgement deadline and do that on a per subscription. So if you know that on average it takes you 15 seconds to process a message in your work queue, then you might set your `ackDeadline` to 20 seconds. This will ensure it doesn't try to redeliver the messages.

## Subscribers can work as a team or separately



Subscribers can work as a individually or as a group. If we have just one subscriber, it is going to get every message delivered through that subscription. However, you can set up worker pools by having multiple subscribers sharing the same subscription.

In this case, it is going to distribute the message, so one and three go to Subscription 1, and two goes to Subscription 2. And it is just random based on when it pulls from messages throughout the day

In the case of a Push subscription, you only have one web inpoint, so you will only have one subscriber typically. But, that one subscriber could be a app engine standard app, or cloud run container image, which autoscales. So, it is one web endpoint, but it can have autoscale workers behind the scenes. And that is actually a very good pattern.

## Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

[Create topic](#)



Let's look at a little bit of code now. This example is using the Client Library for Pub/Sub. If we want to publish the message, we first create the topic.

## Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

[Create topic](#)

```
gcloud pubsub topics publish sandiego "hello"
```

[Publish to topic](#)

Then, we can publish the topic on the command line.

## Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego "hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

Create a client

Set topic name

Message

Send attribute



More commonly, it will be done in code. Here we get a PublisherClient, create topic, and publish the message.

## Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego "hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

Create a client

Set topic name

Message

Send attribute



Notice the `b`, right here in front of 'My first message!' This is because Pub/Sub just sends raw bytes. This means that you aren't restricted to just text. You can send other data, like images if you wanted to. The limit is 10 MB.



# Publishing with Cloud Pub/Sub

```
gcloud pubsub topics create sandiego
```

Create topic

```
gcloud pubsub topics publish sandiego "hello"
```

Publish to topic

```
import os
from google.cloud import pubsub_v1

publisher = pubsub_v1.PublisherClient()

topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)

publisher.create_topic(topic_name)
publisher.publish(topic_name, b'My first message!', author='dylan')
```

Python

Create a client

Set topic name

Message

Send attribute



There are also extra attributes that you can include in messages. In this example, you see author equals dylan. Pub/Sub will keep track of those attributes to allow your downstream systems to get metadata about your messages without having to put it all in the message and parse it out. So instead of serializing and de-serializing, it will just keep track of those key value pairs.

Some of them have special meaning. We will see some of those shortly.

## Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name
)

def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

Python

Create a client

Select  
topic  
name

Set subscription  
name

callback when  
message received

Push method  
Callback function



To subscribe with Pub/Sub using the Push method, the code is similar. Select the topic.

## Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name

def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

Python

Create a client

Select  
topic  
name

Set subscription  
name

callback when  
message received

Push method

Callback function



Name the subscription.

## Subscribing with Cloud Pub/Sub using async pull

```
import os
from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()
topic_name = 'projects/{project_id}/topics/{topic}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    topic='MY_TOPIC_NAME',
)
subscription_name = 'projects/{project_id}/subscriptions/{sub}'.format(
    project_id=os.getenv('GOOGLE_CLOUD_PROJECT'),
    sub='MY_SUBSCRIPTION_NAME',
)
subscriber.create_subscription(
    name=subscription_name, topic=topic_name
)

def callback(message):
    print(message.data)
    message.ack()

future = subscriber.subscribe(subscription_name, callback)
```

Python

Create a client

Select  
topic  
name

Set subscription  
name

callback when  
message received

Push method  
Callback function



This is a push subscription, so we will define a callback.

## Subscribing with Cloud Pub/Sub using synchronous pull

```
gcloud pubsub subscriptions create --topic sandiego mySub1
```

Create subscription

```
gcloud pubsub subscriptions pull --auto-ack mySub1
```

Pull subscription

```
import time

from google.cloud import pubsub_v1

subscriber = pubsub_v1.SubscriberClient()

subscription_path = subscriber.subscription_path(project_id, subscription_name)

NUM_MESSAGES = 2
ACK_DEADLINE = 30
SLEEP_TIME = 10

# The subscriber pulls a specific number of messages.
response = subscriber.pull(subscription_path, max_messages=NUM_MESSAGES)
```

Set subscription name

Create a client

`'projects/{project_id}/subscriptions/{subscription_name}'` ← subscription\_path format


Subscriber is non-blocking

Keep the main thread from exiting to allow it to process messages synchronously



When you are doing a pull subscription, it looks like this. You can pull messages from the command line. You will see this in the lab. By default it will just churn one message, the latest message, but there is a dash-dash limit you can set. Maybe you want 10 messages at a time, you can try that in the lab.

By default, the Publisher batches messages; turn this off if you desire lower latency



Batching messages: throughput versus latency

You can also batch publish messages. This is just prevents the overhead of the call for individual messages on the publisher side. This allows the publisher to wait and send 10, or 50, at a time. This increases efficiency. However, if you are waiting for 50 messages, this means the first one now has latency associated with it. So, it is a trade off in your system. What do you want to optimize? But, in any case, even if you batch publish, they still get delivered one at a time to your subscribers.

We will practice this technique in the lab.

## Changing the batch settings in Cloud Pub/Sub

```
from google.cloud import pubsub
from google.cloud.pubsub import types

client = pubsub.PublisherClient(
    batch_settings=BatchSettings(max_messages=500),
)
```

Python

[Change batch setting](#)



Here is how you would set the batch in python code.

## Pub/Sub: latency, out-of-order, duplication will happen

- Latency -- no guarantees
- Messages can be delivered in any order, especially with large backlog
- Duplication may happen



Just because data can be delivered out of order and in multiple instances, does not mean you have to process the data that way.

You could write an application that handles out-of-order and replicated messages. This is different than a true queueing system. In general, they will be delivered in order, but you can't rely on that with Pub/Sub. This is because, again, one of the compromises made for scalability especially since it is a global service. We have a mesh network, so a message might take another route and if it happens to be a slower route, you could have an earlier message arriving later.

So, for example, you wouldn't use this to implement a chat app because it would be awkward when messages arrive out of order. Therefore, we will handle ordering using other techniques.

Finally, we need to be ready for duplication.



## Cloud Pub/sub with Dataflow: Exactly once, ordered processing




Cloud Pub/Sub delivers at least once



Cloud Dataflow: Deduplicate, order, and window



Separation of concerns → scale




We are going to use Dataflow in conjunction with Pub/Sub to solve these problems. Dataflow will deduplicate messages based on the message ID because in Pub/Sub, if a message is delivered twice, it will have the same ID in both cases. BigQuery also has limited capabilities to deduplicate.

Dataflow will not be able to order in the sense of providing exact sequential order of when messages were published. However, it will help us deal with late data.

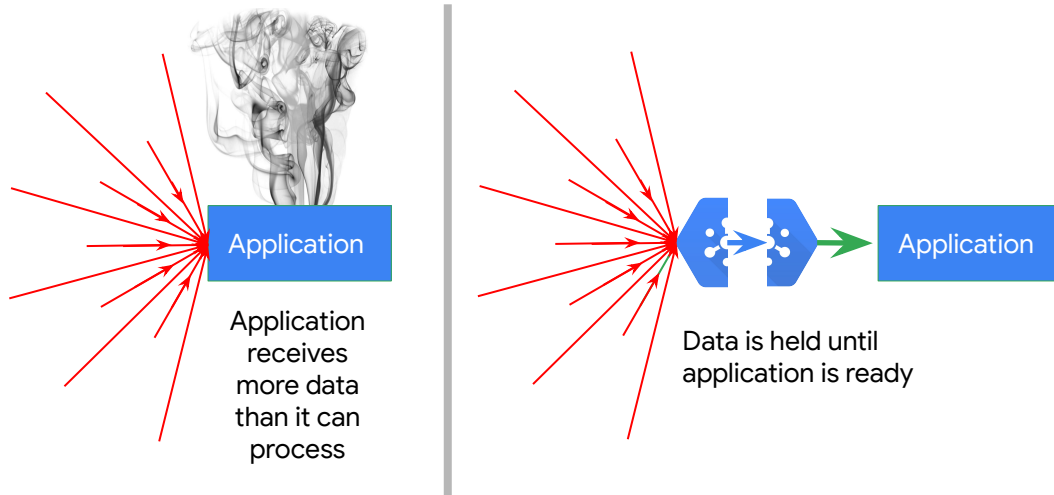
Using Pub/Sub and Dataflow together, allows us to get a scale that wouldn't be possible otherwise.

## Use Cloud Pub/Sub for streaming resilience



Pub/Sub is also going to help us with streaming resilience, or buffering. What happens if your systems get overloaded with large volumes of transactions, like Black Friday? What you really need is some sort of buffer or backlog so that you can feed messages only as fast as the systems are able to process them. Pub/Sub has this as a built in capability.

## Use Cloud Pub/Sub for streaming resilience



Let's recap this then. When you look at the example on this slide, in the example on the left, an overload of arriving data causes a traffic spike. This overdrives the resources of the application as illustrated by the smoke. One solution to this problem is to size the application to handle the highest traffic spike plus some additional capacity as a safety buffer. This is not only wasteful of resources which must be retained at top capacity even when not being used, but it provides a recipe for a distributed denial of service attack by creating an upper limit at which the application will cease to behave normally and will exhibit non-deterministic behavior.

The solution on the right uses Cloud Pub/Sub as an intermediary, receiving and holding data until the application has resources to handle it, either through processing the backlog of work, or by autoscaling to meet the demand.

<https://pixabay.com/photos/smoke-background-artwork-swirl-69124/>

## Security, monitoring, and logging for Pub/Sub

### Cloud Logging metrics

pubsub\_topic  
pubsub\_subscription



### Cloud Audit Logs

Admin = on  
Data Access = off

### Authentication

Service accounts  
User accounts



### Access control

Pub/Sub roles per topic



Cloud Audit Logs maintains three audit logs for each Google Cloud project, folder, and organization: **Admin Activity**, **Data Access**, and **System Event**.

Admin Activity audit logs contain log entries for API calls or other administrative actions that modify the configuration or metadata of resources.

Data Access audit logs contain API calls that read the configuration or metadata of resources, as well as user-driven API calls that create, modify, or read user-provided resource data.

Admin Activity audit logs are always written; you can't configure or disable them.

Data Access audit logs are disabled by default because they can be quite large; they must be explicitly enabled to be written.

Pub/Sub reports metrics to Cloud Logging, including `pubsub_topic` and `pubsub_subscription` which you can monitor against service quota utilization in a dashboard, and for setting notifications and alerts.

Authentication is provided by service accounts. You can also directly authenticate users by their user accounts and their identity is reported in audit logs. But user account authentication is not recommended. Access control is provided by Cloud IAM.



---

Publish Streaming Data  
into Pub/Sub

Now, let's practice publishing streaming data into Pub/Sub.

# Lab Objectives

---

Create a Pub/Sub topic and subscription

Simulate your traffic sensor data into Pub/Sub



What we are trying to do in this lab is to create a Pub/Sub topic and subscription and simulate San Francisco traffic data into Pub/Sub. This is the first of four parts we will do.