



High-Throughput  
BigQuery and Bigtable  
Streaming Features

# Agenda

---

Processing Streaming Data

Cloud Pub/Sub

Cloud Dataflow Streaming  
Features

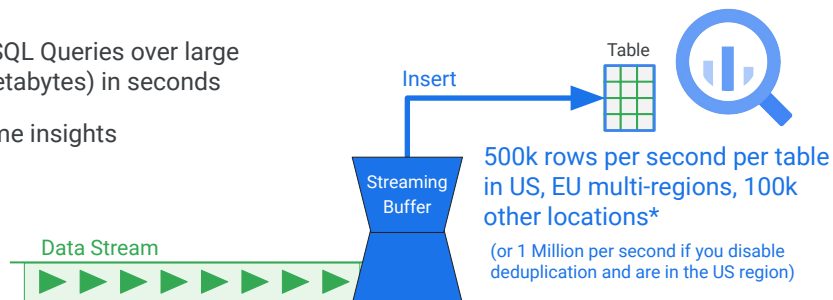
BigQuery and Bigtable Streaming  
Features

Advanced BigQuery Functionality



## BigQuery allows you to stream records into a table; query results incorporate latest data

- Interactive SQL Queries over large datasets (petabytes) in seconds
- Near-real-time insights



Note: Unlike load jobs, there is a cost for streaming inserts (see [quota and limits](#))



Data availability and consistency are considerations. Candidates for streaming are analysis or applications that are tolerant of late or missing data or data arriving out of order or duplicated. The stream can pass through other services introducing additional latency and the possibility of errors.

Since streaming data is unbounded, you need to consider the streaming quotas. There is both a daily limit and a concurrent rate limit. You can find more information about these in the online documentation.

<https://cloud.google.com/bigquery/streaming-data-into-bigquery>

You can disable best effort de-duplication by not populating the `insertId` field for each row inserted. When you do not populate `insertId`, you get much higher streaming ingest quotas for the US region. (1 Million per second vs 500,000 inserts per second)

See:

<https://cloud.google.com/blog/products/data-analytics/whats-happening-bigquery-adding-speed-and-flexibility-10x-streaming-quota-cloud-federation-and-more>

### Question:

When should you ingest a stream of data rather than use a batch approach to load data?

**Answer:**

When the immediate availability of the data is a solution requirement.

**Reason:**

In most cases loading batch data is not charged. Streaming is charged. So use batch loading or repeated batch loading rather than streaming unless that is a requirement of the application.

# Insert streaming data into a BigQuery table

```
export GOOGLE_APPLICATION_CREDENTIALS="/home/user/Downloads/[FILE_NAME].json"
```

## Credentials

The service must have  
Cloud IAM permissions  
set in the Web UI

```
pip install google-cloud-bigquery
```

## Install API

```
from google.cloud import bigquery
client = bigquery.Client(project='PROJECT_ID')

dataset_ref = bigquery_client.dataset('my_dataset_id')
table_ref = dataset_ref.table('my_table_id')
table = bigquery_client.get_table(table_ref)

# read data from Cloud Pub/Sub and place into row format
# static example customer orders in units:
rows_to_insert = [
    (u'customer 1', 5),
    (u'customer 2', 17),
]
errors = bigquery_client.insert_rows(table, rows_to_insert)
```

## Python

## Create a client

## Access dataset and table

## Perform insert



Here is an example of the code used to insert streaming data into a BigQuery table.

## Review streaming data in BigQuery

Query editor										
<pre>1 select * from cloud-training-demos.demos.current_conditions;</pre>										
<div><div>Run</div><div>Save query</div><div>Save view</div><div>Schedule query</div><div>More</div></div>										
Query results										
Query complete (1.2 sec elapsed, 14.3 MB processed)										
<div>Job informationResultsJSONExecution details</div>										
Row	timestamp	latitude	longitude	highway	direction	lane	speed	sensorid		
1	2008-11-01 11:55:00 UTC	33.191415	-117.363042	5	N	4	10.5	33.191415,-117.363042,5,N,4		
2	2008-11-01 11:55:00 UTC	33.191415	-117.363042	5	N	4	10.5	33.191415,-117.363042,5,N,4		
3	2008-11-01 09:35:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4		
4	2008-11-01 12:30:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4		
5	2008-11-01 09:40:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4		
6	2008-11-01 10:55:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4		



After streaming into a BigQuery table has been initiated, you can review the data in BigQuery by querying the table receiving the streaming data.

## Want to visualize insights? Explore Google Data Studio insights right from within BigQuery

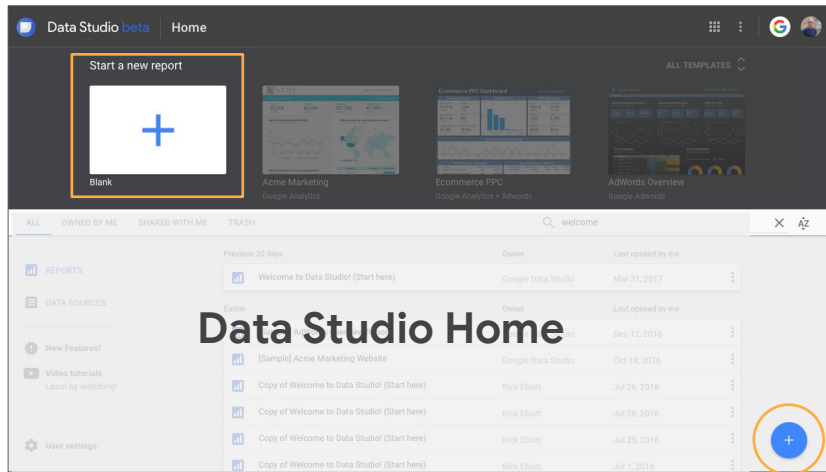
The screenshot shows the Google BigQuery interface. At the top is the 'Query editor' with a SQL query: `select * from cloud-training-demos.demos.current_conditions;`. Below the editor are buttons for 'Run', 'Save query', 'Save view', 'Schedule query', and 'More'. The 'Query results' section shows 'Query complete (1.2 sec elapsed, 14.3 MB processed)'. Below this are tabs for 'Job information', 'Results' (selected), 'JSON', and 'Execution details'. The 'Results' tab displays a table with 6 rows and 10 columns: Row, timestamp, latitude, longitude, highway, direction, lane, speed, sensorid, and an empty column. The 'EXPLORE DATA' button is highlighted with a yellow box.

Row	timestamp	latitude	longitude	highway	direction	lane	speed	sensorid	
1	2008-11-01 11:55:00 UTC	33.191415	-117.363042	5	N	4	10.5	33.191415,-117.363042,5,N,4	
2	2008-11-01 11:55:00 UTC	33.191415	-117.363042	5	N	4	10.5	33.191415,-117.363042,5,N,4	
3	2008-11-01 09:35:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4	
4	2008-11-01 12:30:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4	
5	2008-11-01 09:40:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4	
6	2008-11-01 10:55:00 UTC	33.191415	-117.363042	5	N	4	11.0	33.191415,-117.363042,5,N,4	



When working with data in BigQuery, including streaming data, you can use Data Studio to explore the data further. After you execute a query, you can choose Data Studio from the Explore Data options to immediately start creating visualizations as part of a dashboard.

## Create new reports in the Data Studio UI



Here's a quick walkthrough of the Data Studio UI. Keep in mind the team is continually updating the product so refer to their documentation and examples for additional practice.

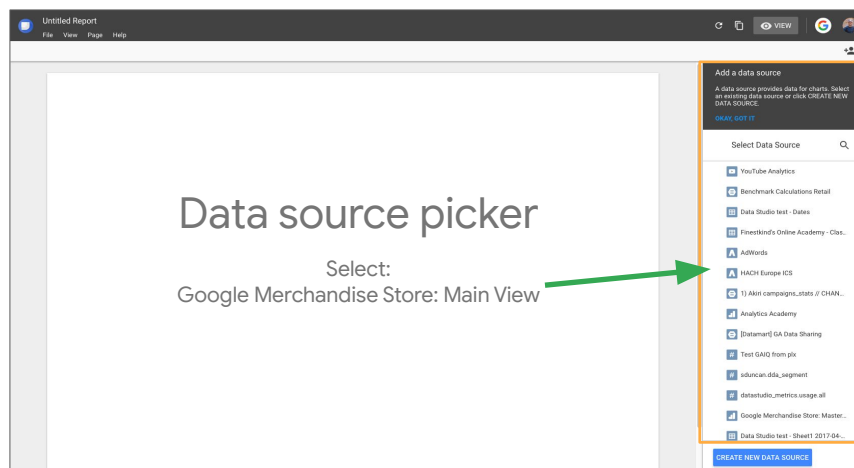
The home page shows the dashboards and data sources you have access to. It's an important distinction -- connected data sources can feed into dashboards but just because someone has access to your dashboard doesn't mean they have permission to view the data presented (because that could be controlled in BigQuery or your GCP project).

Anyway, there two ways to create a new report from scratch: from the templates panel on top

Or from the button in the lower right.



## Select data sources to build your visualizations




The first thing you need to do is tell Data Studio where your data is coming from. That is known as the "data source"

A Data Studio report can have any number of data sources, but we'll just start with one.

The data source picker shows all the data sources you have access to.

## Connect to multiple different types of data sources

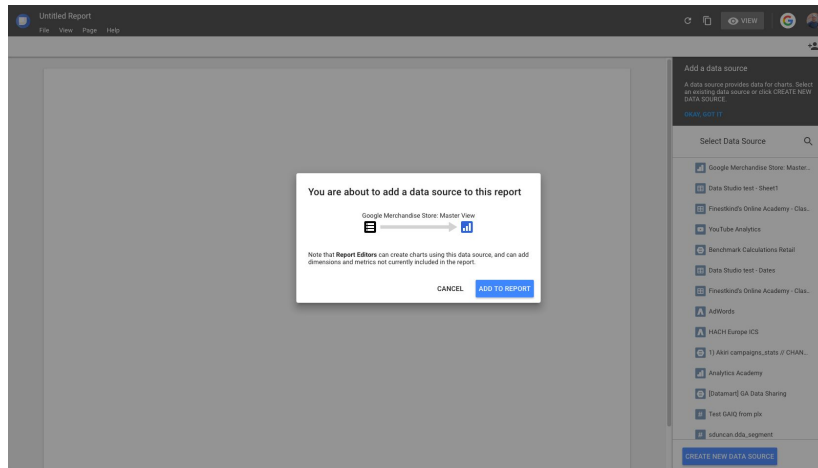


 Google Cloud

Note that you can have any or all of these data sources in a single Data Studio report.

Since Data Studio reports and datasets can be shared, you should be aware of the ramifications of adding a data source.

## Add the data source to your report



When you add a data source to a report, other people who can view the report can potentially see all the data in that data source if you share that source with them. And anyone who can edit the report can use all the fields from any added data sources to create new charts with them.

You can click **ADD TO REPORT** to connect the data source and then you are ready to start visualizing.



---

## Streaming Analytics and Dashboards

### Objectives

- Connect to a BigQuery data source from Google Data Studio
- Create reports and charts to visualize BigQuery data

In this lab, we will connect to BigQuery data source from Google Data Studio and create reports and charts to visualize the BigQuery data.

# Cloud Bigtable

Cloud  
Bigtable



Qualities that Cloud Bigtable contributes to Data Engineering solutions:

NoSQL Queries over large datasets (petabytes) in milliseconds

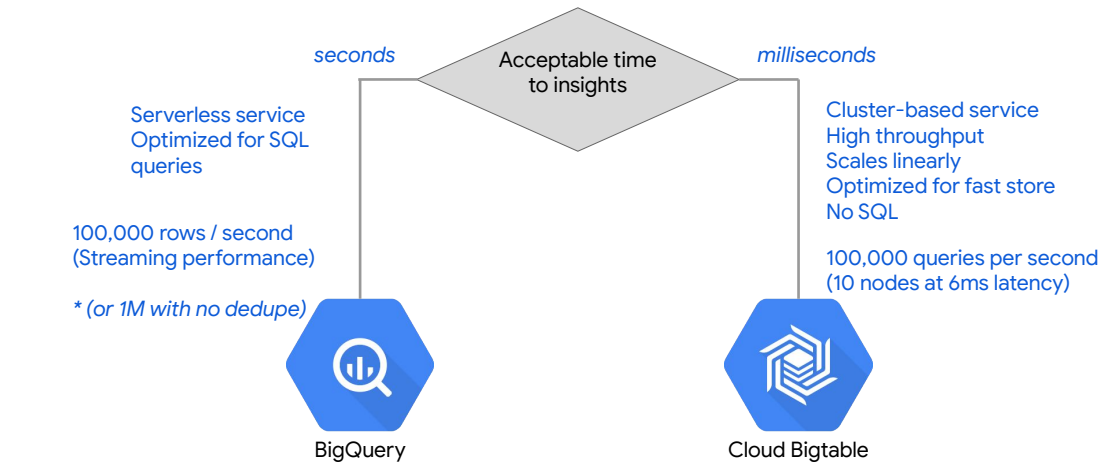
Very fast for specific cases



To use Cloud Bigtable effectively you have to know a lot about your data and how it will be queried up-front. A lot of the optimizations happen before you load data into Cloud Bigtable.


<https://pixabay.com/illustrations/chess-black-and-white-pieces-3413429/>

# How to choose between Cloud Bigtable and BigQuery



Cloud Bigtable is ideal for applications that need very high throughput and scalability for non-structured key/value data, where each value is typically no larger than 10 MB. Not good for highly structured data, transactional data, small data (less than 1 TB), and anything requiring SQL Queries and SQL-like joins.

## Consider Cloud Bigtable for these requirements



Here are a few examples of data engineering requirements that have been solved using Cloud Bigtable.


Machine learning algorithms frequently have many or all of these requirements. Application that use marketing data, such as purchase histories or customer preferences. Applications that use financial data such as transaction histories, stock prices, or currency exchange rates. Internet of Things, IoT data, such as usage reports from meters, sensors, or devices. Time-series data, such as resource consumption like CPU and memory usage over time for multiple servers.

The most common use of Cloud Bigtable is...

Productionize a real-time lookup as part of an application, where speed and efficiency are desired beyond that of other databases.



## How does Cloud Bigtable work?



 Google Cloud


Cloud Bigtable stores data in a file system called Colossus. Colossus also contains data structures called Tablets that are used to identify and manage the data. And metadata about the Tablets is what is stored on the VMs in the Bigtable cluster itself.

This design provides amazing qualities to Cloud Bigtable. It has three levels of operation. It can manipulate the actual data. It can manipulate the Tablets that point to and describe the data. Or it can manipulate the metadata that points to the Tablets. Rebalancing tablets from one node to another is very fast, because only the pointers are updated.

Cloud Bigtable is a learning system. It detects "hot spots" where a lot of activity is going through a single Tablet and splits the Tablet in two. It can also rebalance the processing by moving the pointer to a Tablet to a different VM in the cluster. So its best use case is with big data -- above 300 GB -- and very fast access but constant use over a longer period of time. This gives Cloud Bigtable a chance to learn about the traffic pattern and rebalance the Tablets and the processing.

When a node is lost in the cluster, no data is lost. And recovery is fast because only the metadata needs to be copied to the replacement node. Colossus provides better durability than the default 3 replicas provided by HDFS.

## Cloud Bigtable design idea is "simplify for speed"



The Row Key is the index.

And you get only one.



Cloud Bigtable stores data in tables. And to begin with, it is just a table with rows and columns. However, unlike other table-based data systems like spreadsheets and SQL databases, Cloud Bigtable has only one index.


That index is called the Row Key. There are no alternate indexes or secondary indexes. And when data is entered, it is organized lexicographically by the Row Key.

The design principle of Cloud Bigtable is speed through simplification. If you take a traditional table, and simplify the controls and operations you allow yourself to perform on it then you can optimize for those specific tasks.

It is the same idea behind RISC (Reduced Instruction Set Computing). Simplify the operations. And when you don't have to account for variations, you can make those that remain very fast.

In Cloud Bigtable, the first thing we must abandon in our design is SQL. This is a standard of all the operations a database can perform. And to speed things up we will drop most of them and build up from a minimal set of operations. That is why Cloud Bigtable is called a NoSQL database.

But speed depends on your data and Row Key



The green items are the results you want to produce from the query. In the best case you are going to scan the Row Key one time, from the top-down. And you will find all the data you want to retrieve in adjacent and contiguous rows. You might have to skip some rows. But the query takes a single scan through the index from top-down to collect the result set.

The second instance is sorting. You are still only looking at the Row Key. In this case the yellow line contains data that you want, but it is out of order. You can collect the data in a single scan, but the solution set will be disorderly. So you have to take the extra step of sorting the intermediate results to get the final results. Now think about this. What does the additional sorting operation do to timing? It introduces a couple of variables. If the solution set is only a few rows, then the sorting operation will be quick. But if the solution set is huge, the sorting will take more time. The size of the solution set becomes a factor in timing. The orderliness of the original data is another factor. If most of the rows are already in order, there will be less manipulation required than if there are many rows out of order. The orderliness of the original data becomes a factor in timing. So introducing sorting means that the time it takes to produce the result is much more variable than scanning.

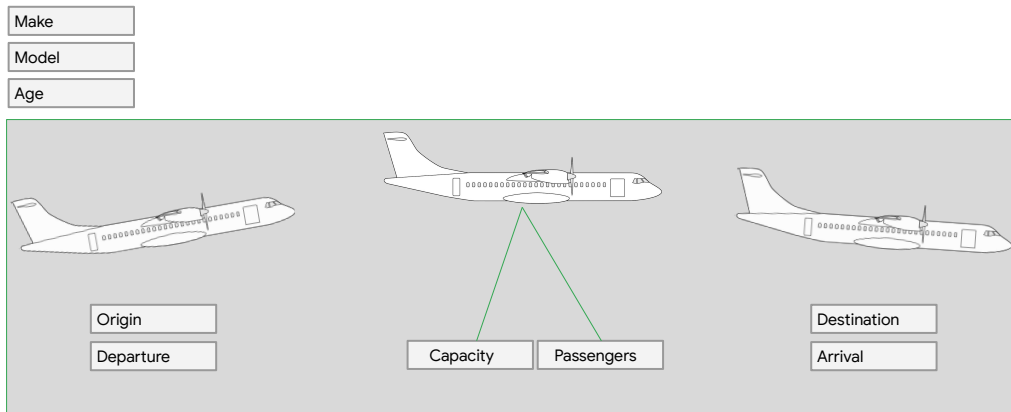
The third instance is searching. In this case, one of the columns contains critical data. You can't tell whether a row is a member of the solution set or not without examining the data contained in the critical column. The Row Key is no longer sufficient. So now you are bouncing back and forth between Row Key and column contents. There are many approaches to searching. You could divide it up into multiple steps, one scan

through the Row Keys and subsequent scans through the columns, and then perhaps a final sort to get the data in the order you want. And it gets much more complicated if there are multiple columns containing critical information. And it gets more complicated if the conditions of solution set membership involve logic such as a value in one column AND a value in another column, or a value in one column OR a value in another column. However, any algorithm or strategy you use to produce the result is going to be slower and more variable than scanning or sorting.

What is the lesson from this exploration? That to get the best performance with the design of the Cloud Bigtable service, you need to get your data in order first, if possible, and you need to select or construct a Row Key that minimizes sorting and searching and turns your most common queries into scans.

Not all data and not all queries are good use cases for the efficiency that the Cloud Bigtable service offers. But when it is a good match, Cloud Bigtable is so consistently fast that it is magical.

## Flights of the world: Reviewing the data



Each entry records the occurrence of one flight.

The data include city of origin and the date and time of departure, and destination city and date and time of arrival.

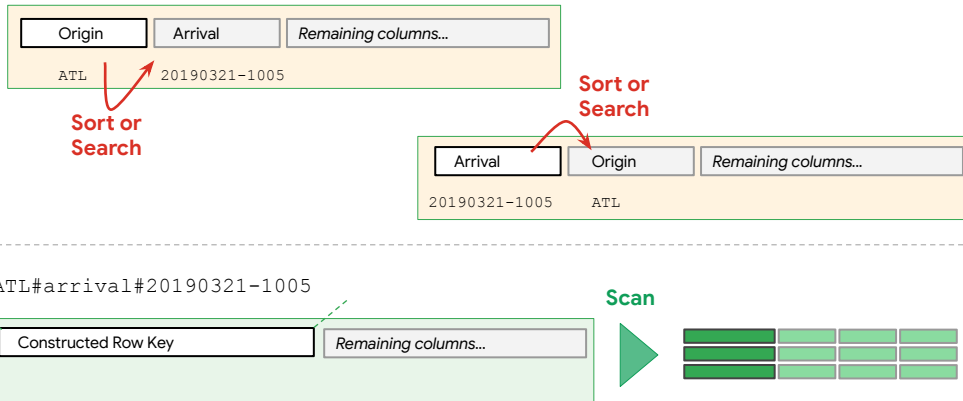
Each airplane has a maximum capacity, and related to this is the number of passengers that were actually aboard each flight.

Finally, there is information about the aircraft itself, including the manufacturer, called the make, the model number, and the current age of the aircraft at the time of the flight.

<https://pixabay.com/vectors/atr-72-aircraft-sideview-drawing-884214/>

## What is the best Row Key?

Query: All flights originating in Atlanta and arriving between March 21st and 29th



In this example, the Row Key will be defined for the most common use case. The Query is to find all flights originating from the Atlanta airport and arriving between March 21st and 29th. The airport where the flight originates is in the Origin field. And the date when the aircraft landed is listed in the Arrival field.

If you use Origin as the Row Key, you will be able to pull out all flights from Atlanta -- but the Arrival field will not necessarily be in order. So that means searching through the column to produce the solution set.

If you use the Arrival field as the Row Key, it will be easy to pull out all flights between March 21st and 29th, but the airport of origin won't be organized. So you will be searching through the arrival column to produce the solution set.

In the third example, a Row Key has been constructed from information extracted from the Origin field and the Arrival field -- creating a constructed Row Key. Because the data is organized lexicographically by the Row Key, all the Atlanta flights will appear in a group, and sorted by date of arrival. Using this Row Key you can generate the solution set with only a scan.

In this example, the data was transformed when it arrived. So constructing a Row Key during the transformation process is straightforward.

## Cloud Bigtable schema organization



### Column Families

Row Key	Flight_Information					Aircraft_Information			
	Origin	Destination	Departure	Arrival	Passengers	Capacity	Make	Model	Age
ATL#arrival#20190321-1121	ATL	LON	20190321-0311	20190321-1121	158	162	B	737	18
ATL#arrival#20190321-1201	ATL	MEX	20190321-0821	20190321-1201	187	189	B	737	8
ATL#arrival#20190321-1716	ATL	YVR	20190321-1014	20190321-1716	201	259	B	757	23



Cloud Bigtable also provide Column Families. By accessing the Column Family, you can pull some of the data you need without pulling all of the data from the row or having to search for it and assemble it. This makes access more efficient.

<https://cloud.google.com/bigtable/docs/schema-design>

<https://pixabay.com/vectors/airplane-jet-aircraft-plane-flight-309503/>

Queries that use the row key, a row prefix, or a row range are the most efficient

Query: Current arrival delay for flights from Atlanta

1

**ROW KEY BASED ON ATLANTA ARRIVALS**  
E.G. ORIGIN#arrival  
( ATL#arrival#20190321-1005 )

Puts latest flights at bottom of table

2

**REVERSE TIMESTAMP TO THE ROWKEY**  
E.G. ORIGIN#arrival#RTS  
( ATL#arrival#12345678 )

Puts latest flights at top of table



most common query is for current arrival delay in Atlanta. That will involve averaging flight delays over the last 30 minutes. Hence, origin#arrival. We want this at the top of the table, hence RTS.



Use reverse timestamps when your most common query is for the latest values.

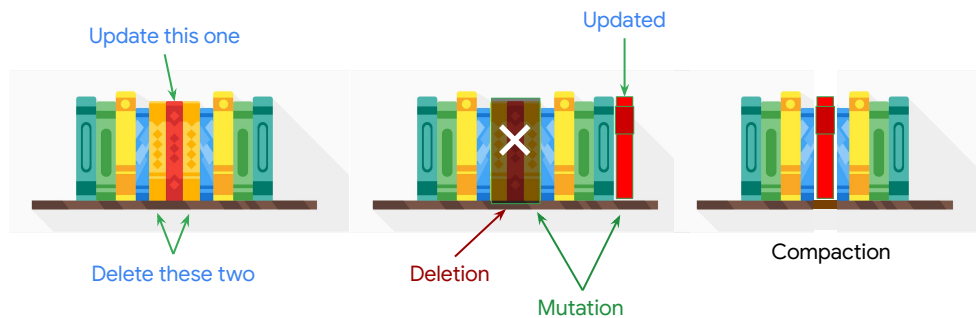
Query: Current arrival delay for flights from Atlanta

```
// key is ORIGIN#arrival#REVTS
String key = info.getORIGIN() //
    + "#arrival" //
    + "#" + (Long.MAX_VALUE - ts.getMillis()); // reverse timestamp
```



You can reverse timestamps by subtracting the timestamp from your programming language's maximum value for long integers (such as Java's `java.lang.Long.MAX_VALUE`, for example `LONG_MAX - timestamp.millisecondsSinceEpoch()`). By reversing the timestamp, you can design a row key where the most recent event appears at the start of the table instead of the end. As a result, you can get the N most recent events simply by retrieving the first N rows of the table.

## What happens when data in Cloud Bigtable is changed?



When you delete data, the row is marked for deletion and skipped during subsequent processing. It is not immediately removed.

If you make a change to data, the new row is appended sequentially to the end of the table, and the previous version is marked for deletion. So both rows exist for a period of time. Periodically, Cloud Bigtable compacts the table, removing rows marked for deletion and reorganizing the data for read and write efficiency.

<https://pixabay.com/vectors/book-rack-shelf-furniture-design-2943383/>

## Optimizing data organization for performance



Group related data for more efficient reads

Example row key:

DehliIndia#2019031411841

Use column families



Distribute data evenly for more efficient writes



Place identical values in the same row or adjoining rows for more efficient compression

Use row keys to organize identical data



Distributing the writes across nodes provides the best write performance. One way to accomplish this is by choosing row keys that are randomly distributed.

However, choosing a row key that groups related rows so they are adjacent makes it much more efficient to read multiple rows at one time.

In our airline example, if we were collecting weather data from the airport cities, we might construct a key consisting of a hash of the city name along with a timestamp. The example row key shown would enable pulling all the data for Dehli, India as a contiguous range of rows.

Whenever there are rows containing multiple column values that are related, it is a good idea to group them into a column family. Some NoSQL databases suffer performance degradation if there are too many column families. Cloud Bigtable can handle up to 100 column families without losing performance. And it is much more efficient to retrieve data from one or more column families than retrieving all of the data in a row.

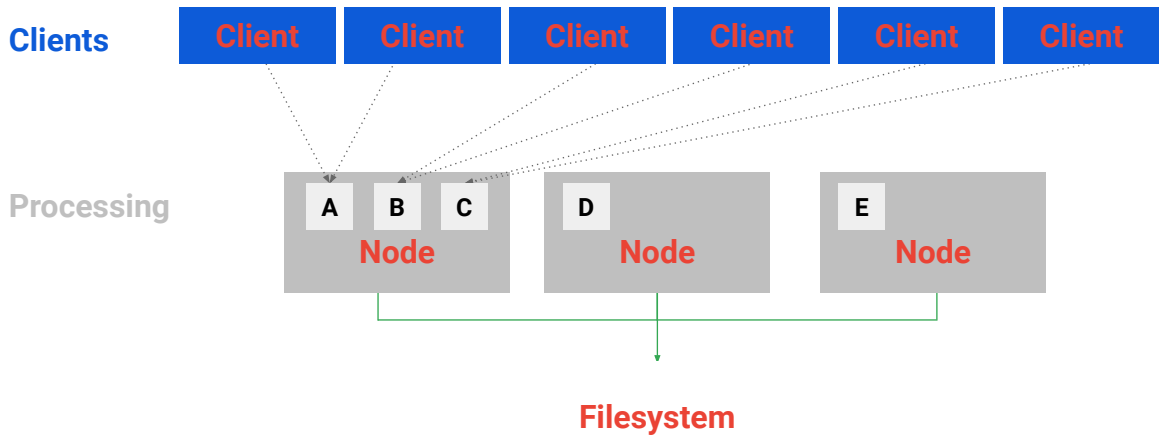
There are currently no configuration settings in Cloud Bigtable for compression. However, random data cannot be compressed as efficiently as organized data. Compression works best if identical values are near each other, either in the same row or in adjoining rows. If you arrange your row keys so that rows with identical data are adjacent, the data can be compressed more efficiently.

<https://pixabay.com/photos/glasses-reading-glasses-302251/>

<https://pixabay.com/photos/colored-pencils-pens-colorful-paint-4030202/>

<https://pixabay.com/photos/band-orchestra-instrument-music-1492367/>

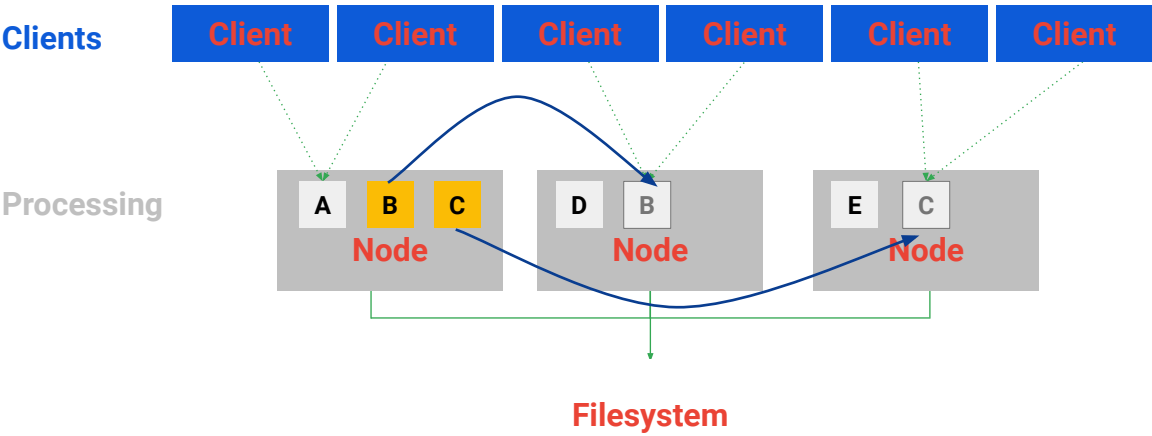
## Cloud Bigtable self-improves by learning access patterns...




### Notes:

A,B,C,D,E are not data but rather pointers/references and cache....which is why rebalancing is not time consuming...we're just moving pointers. Actual data is in tablets in Colossus FS.

...and rebalances data accordingly



### Rebalance strategy: distribute reads



### Notes:

Cloud Bigtable tries to distribute reads and writes equally across all Cloud Bigtable nodes.

With a well-designed schema, reads and writes should be distributed fairly evenly across an entire table and cluster. However, in some cases, it is inevitable that certain rows will be accessed more frequently than others.

In these cases, Cloud Bigtable will redistribute tablets so that reads are spread evenly across nodes in the cluster.

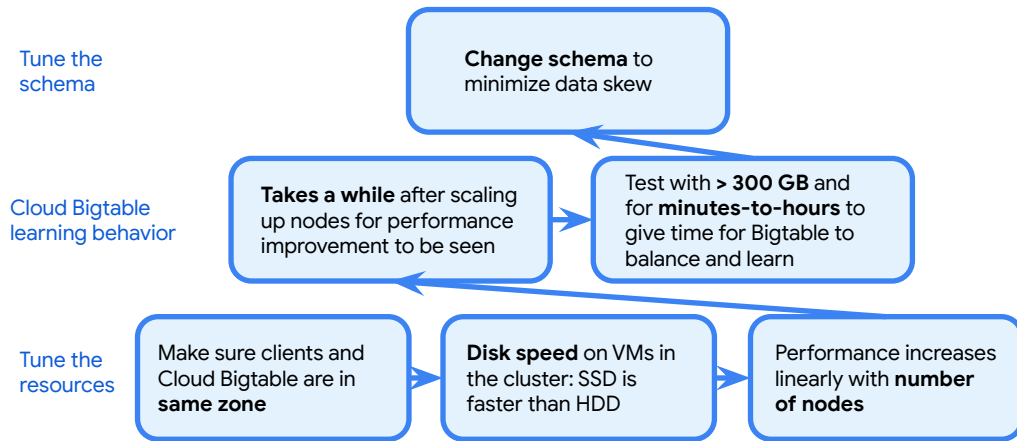
Note that ensuring an even distribution of reads has taken priority over evenly distributing storage across the cluster.

# Optimizing Cloud Bigtable Performance





# Optimizing Cloud Bigtable Performance



There are several factors that can result in slower performance:

The table's schema is not designed correctly

It's essential to design a schema that allows reads and writes to be evenly distributed across the Cloud Bigtable cluster. Otherwise, individual nodes can get overloaded, slowing performance.

The workload isn't appropriate for Cloud Bigtable

Testing with a small amount (< 300 GB) of data, or for a very short period of time (seconds rather than minutes or hours), Cloud Bigtable won't be able to properly optimize your data. It needs time to learn your access patterns, and it needs large enough shards of data to make use of all of the nodes in your cluster.

The Cloud Bigtable cluster doesn't have enough nodes

Typically, performance increases linearly with the number of nodes in a cluster.

Adding more nodes can therefore improve performance. Use the monitoring tools to check whether a cluster is overloaded.

The Cloud Bigtable cluster was scaled up very recently

While nodes are available in your cluster almost immediately, Cloud Bigtable can take up to 20 minutes under load to optimally distribute cluster workload across the new nodes.

The Cloud Bigtable cluster uses HDD disks


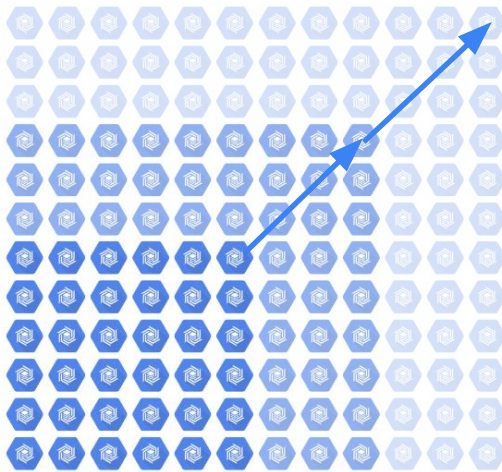
Using HDD disks instead of SSD disks means slower response times and a significantly lower cap on the number of read requests handled per second (500 QPS for HDD disks vs. 10,000 QPS for SSD disks).

There are issues with the network connection


Network issues can reduce throughput and cause reads and writes to take longer than usual. In particular, you'll see issues if your clients are not running in the same zone as your Cloud Bigtable cluster.

Because different workloads can cause performance to vary, you should perform tests with your own workloads to obtain the most accurate benchmarks.

## Throughput can be controlled by node count




## Features for Cloud Bigtable streaming



There are some things you can do that will improve write performance and read performance.

There are some things you can do that are a trade-off where improving write may cost you on read

# Schema design is the primary control for streaming



A higher throughput means more items are processed in a given amount of time.

If you have larger rows, then fewer of them will be processed in the same amount of time.

In general, smaller rows offers higher throughput, and therefore is better for streaming performance.

Cloud Bigtable takes time to process cells within a row.

So if there are fewer cells within a row, it will generally provide better performance than more cells.

Finally, selecting the right row key is critical. Rows are sorted lexicographically.

The goal when optimizing for streaming is to avoid creating hotspots when writing, which would cause Cloud Bigtable to have to split tablets and adjust loads.

To accomplish that, you want the data to be as evenly distributed as possible.

<https://cloud.google.com/bigtable/docs/performance>


reading delay + processing delay = response time

## Use Cloud Bigtable replications to improve availability

### Why perform replication?

- Isolate serving applications from batch reads
- Improve availability
- Provide near-real-time backup
- Ensure your data has a global presence

```
gcloud bigtable clusters create CLUSTER_ID \  
  --instance=INSTANCE_ID \  
  --zone=ZONE \  
  [--num-nodes=NUM_NODES] \  
  [--storage-type=STORAGE_TYPE]
```



Replication for Cloud Bigtable enables you to increase the availability and durability of your data by copying it across multiple regions or multiple zones within the same region. You can also isolate workloads by routing different types of requests to different clusters.

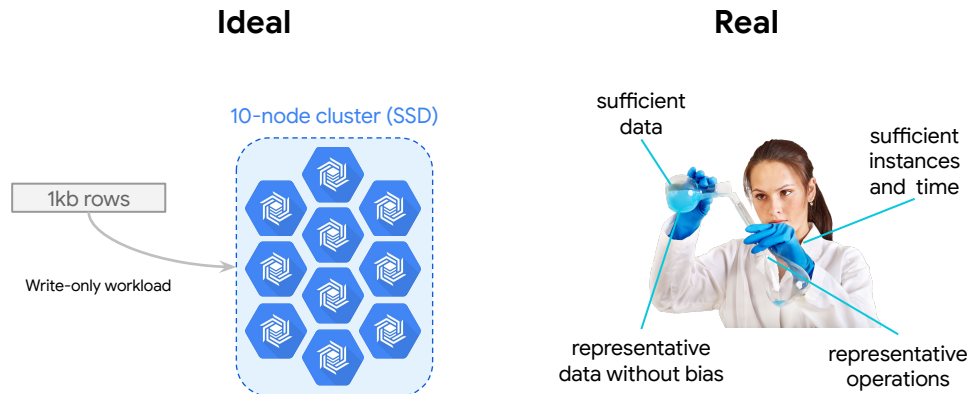
Simple call **gcloud bigtable clusters create** to create a cluster of Bigtable replicas.

If a Cloud Bigtable cluster becomes unresponsive, replication makes it possible for incoming traffic to fail over to another cluster in the same instance. Failovers can be either manual or automatic, depending on the app profile an application is using and how the app profile is configured.

The ability to create multiple clusters in an instance is valuable for performance, as one can be for writing and the replica cluster exclusively for reading. BT also supports automatic failover for HA

<https://cloud.google.com/bigtable/docs/replication-overview>

## Run performance tests carefully for Cloud Bigtable streaming



The generalizations, of isolate the write workload, increase number of nodes, and decrease row size and cell size will not apply in all cases.

In most circumstances, experimentation is the key to defining the best solution.

A performance estimate is given in the documentation online for write-only workloads. Of course, the purpose of writing data is to eventually read it, so the baseline is an ideal case.

At the time of this writing, a 10-node SSD cluster with 1kb rows and a write-only workload can process 10,000 rows per second at a 6ms delay.


This estimate will be affected by average row size, the balance and timing of reads distracting from writes, and other factors.

You will want to run performance tests with your actual data and application code. You need to run the tests on at least 300 GB of data to get valid results.

Also, to get valid results your test needs to perform enough actions over a long enough period of time to give Cloud Bigtable the time and conditions necessary to learn the usage pattern and perform its internal optimizations.

<https://cloud.google.com/bigtable/docs/performance>  
<https://pixabay.com/photos/chemistry-lab-experiment-3005692/>

## Key Visualizer exposes read/write access patterns over time and key space



Key Visualizer is a tool that helps you analyze your Cloud Bigtable usage patterns. It generates visual reports for your tables that break down your usage based on the row keys that you access. Key Visualizer automatically generates hourly and daily scans for every table in your instance that meets at least one of the following criteria:

- During the previous 24 hours, the table contained at least 30 GB of data at some point in time.
- During the previous 24 hours, the average of all reads or all writes was at least 10,000 rows per second.

The core of a Key Visualizer scan is the heatmap, which shows the value of a metric over time, broken down into contiguous ranges of row keys. The x-axis of the heatmap represents time, and the y-axis represents row keys. If the metric had a low value for a group of row keys at a point in time, the metric is "cold," and it appears in a dark color. A high value is "hot," and it appears in a bright color; the highest values appear in white.





---

## Streaming Data Pipelines into Bigtable

### Objectives

- Launch a Dataflow pipeline to read from PubSub and write into Bigtable
- Open an HBase shell to query the Bigtable database