



Advanced BigQuery Functionality

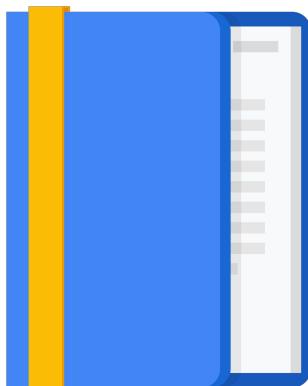
Agenda

Analytic Window Functions

Using With Clauses

GIS Functions

Performance Considerations



Use analytic window functions for advanced analysis

- Standard aggregations
- Navigation functions
- Ranking and numbering functions



Example: The COUNT function (self-explanatory)

SELECT
start_station_name,
end_station_name,
APPROX_QUANTILES(duration, 10)[OFFSET(5)] AS typical_duration,
COUNT(*) AS num_trips
FROM
`bigquery-public-data`.london_bicycles.cycle_hire
GROUP BY
start_station_name,
end_station_name

Query results [SAVE RESULTS](#) ▾ [EXPLORE WITH DATA STUDIO](#)

Query complete (13.5 sec elapsed, 1.5 GB processed)

Job information [Results](#) [JSON](#) [Execution details](#)

Row	start_station_name	end_station_name	typical_duration	num_trips
1	Borough High Street, The Borough	Bell Street, Marylebone	4320	3
2	William IV Street, Strand	Little Brook Green, Brook Green	4500	3
3	Baker Street, Marylebone	George Place Mews, Marylebone	180	95
4	Waterloo Station 2, Waterloo	Westbourne Grove, Bayswater	3240	7
5	Imperial Wharf Station	Upperne Road, West Chelsea	180	94
6	Kennington Road , Vauxhall	Westminster Bridge Road, Elephant & Castle	180	38
7	Whiston Road, Haggerston	Pitfield Street North,Hoxton	180	99
8	Gloucester Street, Pimlico	Rampayne Street, Pimlico	180	330
9	Harrington Square 1, Camden Town	Drummond Street , Euston	180	76



Some other “standard” aggregation functions

- SUM
- AVG
- MIN
- MAX
- BIT_AND
- BIT_OR
- BIT_XOR
- COUNTIF
- LOGICAL_OR
- LOGICAL_AND

https://cloud.google.com/bigquery/docs/reference/standard-sql/aggregate_functions



Example: The LEAD function returns the value of a row n rows ahead of the current row

```
SELECT
  start_date,
  end_date,
  bike_id,
  start_station_id,
  end_station_id,
  LEAD(start_date, 1) OVER(PARTITION BY bike_id ORDER BY start_date ASC ) AS next_rental_start
FROM
  `bigquery-public-data`.london_bicycles.cycle_hire
WHERE
  bike_id = 9
```

Query results  

Query complete (2.0 sec elapsed, 926.1 MB processed)

Job Information [Results](#) [JSON](#) [Execution details](#)

Row	start_date	end_date	bike_id	start_station_id	end_station_id	next_rental_start
1	2015-01-04 14:03:00 UTC	2015-01-04 15:17:00 UTC	9	176	176	2015-01-05 09:04:00 UTC
2	2015-01-05 09:40:00 UTC	2015-01-05 09:22:00 UTC	9	176	106	2015-01-05 18:17:00 UTC
3	2015-01-05 18:17:00 UTC	2015-01-05 18:32:00 UTC	9	106	286	2015-01-06 16:23:00 UTC
4	2015-01-06 16:23:00 UTC	2015-01-06 16:30:00 UTC	9	286	99	2015-01-06 17:08:00 UTC
5	2015-01-06 17:08:00 UTC	2015-01-06 17:14:00 UTC	9	99	49	2015-01-06 17:51:00 UTC
6	2015-01-06 17:51:00 UTC	2015-01-06 18:02:00 UTC	9	49	345	2015-01-06 18:58:00 UTC
7	2015-01-06 18:58:00 UTC	2015-01-06 19:13:00 UTC	9	345	603	2015-01-07 08:30:00 UTC
8	2015-01-07 08:30:00 UTC	2015-01-07 08:48:00 UTC	9	603	112	2015-01-07 16:44:00 UTC
9	2015-01-07 16:44:00 UTC	2015-01-07 16:58:00 UTC	9	465	67	2015-01-07 17:05:00 UTC



Some other navigation functions

- NTH_VALUE
- LAG
- FIRST_VALUE
- LAST_VALUE

https://cloud.google.com/bigquery/docs/reference/standard-sql/navigation_functions



Example: The RANK function returns the integer rank of a value in a group of values

```
WITH
longest_trips AS (
SELECT
start_station_id,
duration,
RANK() OVER(PARTITION BY start_station_id ORDER BY duration DESC) AS nth_longest
FROM
`bigquery-public-data`.london_bicycles.cycle_hire )
SELECT
start_station_id,
ARRAY_AGG(duration
ORDER BY
nth_longest
LIMIT
3)AS durations
FROM
longest_trips
GROUP BY
start_station_id
```

Query results  SAVE RESULTS EXPLORE WITH DATA STUDIO

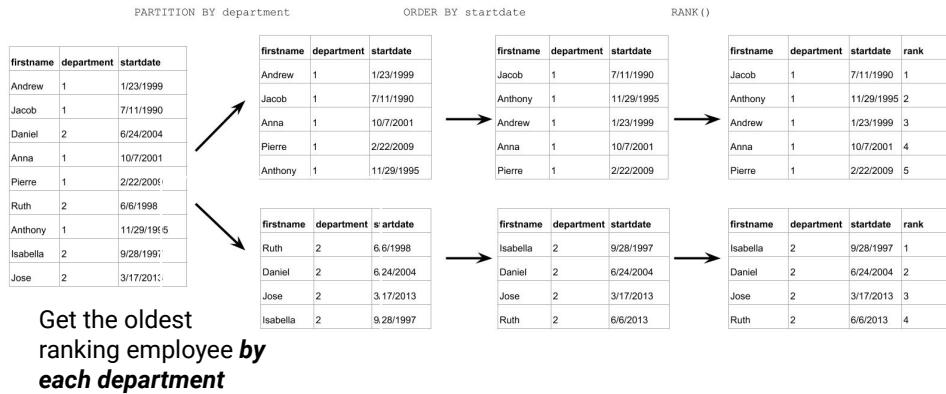
Query complete (7.0 sec elapsed, 370.1 MB processed)

Job information Results  JSON Execution details

Row	start_station_id	durations
1	4	457380 406140 405300
2	11	872400 511680 312420
3	21	2238840 1188000 587340



Example: RANK() function for aggregating over groups of rows



Get the oldest ranking employee by each department

<https://cloud.google.com/bigquery/docs/reference/standard-sql/functions-and-operators#supported-functions>
RANK()

In databases, an analytic function is a function that computes aggregate values over a group of rows. Unlike aggregate functions, which return a single aggregate value for a group of rows, analytic functions return a single value for each row by computing the function over a group of input rows.

Analytic functions are a powerful mechanism for succinctly representing complex analytic operations, and they enable efficient evaluations that otherwise would involve expensive self-JOINS or computation outside the SQL query.

Analytic functions are also called "(analytic) window functions" in the SQL standard and some commercial databases. This is because an analytic function is evaluated over a group of rows, referred to as a window or window frame. In some other databases, they may be referred to as Online Analytical Processing (OLAP) functions.

Example: RANK() function for aggregating over groups of rows

```
SELECT firstname, department, startdate,  
       RANK() OVER ( PARTITION BY department ORDER BY startdate )  
AS rank  
FROM Employees;
```



Option to do demo with IRS dataset:

```
#standardSQL  
# Largest employer per U.S. state per 2015 filing  
WITH employer_per_state AS (  
SELECT  
    ein,  
    name,  
    state,  
    noemployeesw3cnt AS number_of_employees,  
    RANK() OVER (PARTITION BY state ORDER BY noemployeesw3cnt DESC ) AS  
rank  
FROM  
    `bigquery-public-data.irs_990.irs_990_2015`  
JOIN  
    `bigquery-public-data.irs_990.irs_990_ein`  
USING(ein)  
GROUP BY 1,2,3,4 #remove duplicates  
)  
  
# Get the top employer per state and order highest to lowest states  
SELECT *  
FROM employer_per_state  
WHERE rank = 1
```

ORDER BY number_of_employees DESC;

Some other ranking and numbering functions

- CUME_DIST
- DENSE_RANK
- ROW_NUMBER
- PERCENT_RANK

https://cloud.google.com/bigquery/docs/reference/standard-sql/numbering_functions



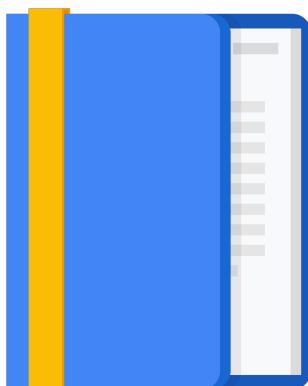
Agenda

Analytic Window Functions

Using With Clauses

GIS Functions

Performance Considerations



Use WITH clauses and subqueries to modularize

```

3 ← WITH
4   # 2015 filings joined with organization details
5   irs_990_2015_ein AS (
6     SELECT *
7     FROM
8       `bigquery-public-data.irs_990.irs_990_2015`
9     JOIN
10      `bigquery-public-data.irs_990.irs_990_ein` USING (ein)
11    ),
12
13   # duplicate EINs in organization details
14   duplicates AS (
15     SELECT
16       ein AS ein,
17       COUNT(ein) AS ein_count
18     FROM
19       irs_990_2015_ein
20     GROUP BY
21       ein
22     HAVING
23       ein_count > 1
24     )
25
26
27   # return results to store in a permanent table
28   SELECT
29     irs_990.ein AS ein,
30     irs_990.name AS name,
31     irs_990.noemployees3cnt AS num_employees,
32     irs_990.grerptspublicuse AS gross_receipts
33     # more fields omitted for brevity
34   FROM irs_990_2015_ein AS irs_990
35   LEFT JOIN duplicates
36   ON
37     irs_990.ein=duplicates.ein
38   WHERE
39     # filter out duplicate records
40     duplicates.ein IS NULL

```



- WITH is simply a named subquery (or Common Table Expression)
- Acts as a temporary table
- Breaks up complex queries
- Chain together multiple subqueries in a single WITH
- You can reference other subqueries in future subqueries

```

#standardSQL
#CTEs
WITH

# 2015 filings joined with organization details
irs_990_2015_ein AS (
SELECT *
FROM
`bigquery-public-data.irs_990.irs_990_2015`
JOIN
`bigquery-public-data.irs_990.irs_990_ein` USING (ein)
),

# duplicate EINs in organization details
duplicates AS (
SELECT
ein AS ein,
COUNT(ein) AS ein_count
FROM
irs_990_2015_ein
GROUP BY
ein
HAVING
ein_count > 1
)

```

```
)  
  
# return results to store in a permanent table  
SELECT  
    irs_990.ein AS ein,  
    irs_990.name AS name,  
    irs_990.noemployeesw3cnt AS num_employees,  
    irs_990.grsrcptspublicuse AS gross_receipts  
    # more fields ommited for brevity  
FROM irs_990_2015_ein AS irs_990  
LEFT JOIN duplicates  
ON  
    irs_990.ein=duplicates.ein  
WHERE  
    # filter out duplicate records  
    duplicates.ein IS NULL  
  
LIMIT 10
```

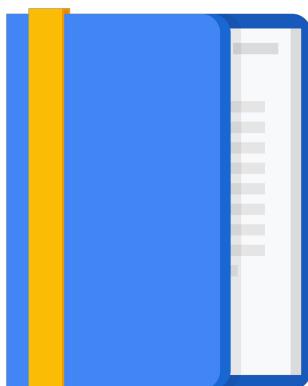
Agenda

Analytic Window Functions

Using With Clauses

GIS Functions

Performance Considerations



BigQuery has built-in GIS functionality

- Example: Can we find the zip codes best served by the New York Citibike system by looking for the number of stations within 1 km of each zip code that have at least 30 bikes?

```
SELECT
    z.zip_code,
    COUNT(*) AS num_stations
FROM
    `bigquery-public-data.new_york_citibike.citibike_stations` AS s,
    `bigquery-public-data.geo_us_boundaries.zip_codes` AS z
WHERE
    ST_DWithin(z.zcta_geom,
        ST_GeogPoint(s.longitude, s.latitude),
        1000) -- 1km
    AND num_bikes_available > 30
GROUP BY
    z.zip_code
ORDER BY
    num_stations DESC
```

Query results [SAVE RESULTS](#) [EXPLORE WITH DATA STUDIO](#)

Query complete (1.1 sec elapsed, 128.3 MB processed)

Job information [Results](#) [JSON](#) [Execution details](#)

Row	zip_code	num_stations
1	10003	21
2	10002	19
3	10026	17
4	10012	16
5	10029	16



Use ST_DWITHIN to check if two locations objects are within some distance

- ST_DWithin(geography_1, geography_2, distance)

 - in meters
- Can lessen resolution of geo coordinates with ST_SnapToGrid
 - ST_SnapToGrid(pt, 0.01) rounds off the latitude and longitude of pt to the second decimal place

```
SELECT
z.zip_code,
COUNT(*) AS num_stations
FROM
`bigquery-public-data.new_york_citibike.citibike_stations` AS s,
`bigquery-public-data.geo_us_boundaries.zip_codes` AS z
WHERE
ST_DWithin(z.zip_code_geom,
ST_GeogPoint(s.longitude, s.latitude),
1000) -- 1km
AND num_bikes_available > 30
GROUP BY
z.zip_code
ORDER BY
num_stations DESC
```



Represent longitude and latitude points as Well Known Text (WKT) using the function ST_GeogPoint

- Queries in BigQuery are much more efficient if geographic data is stored as geography types rather than as primitives

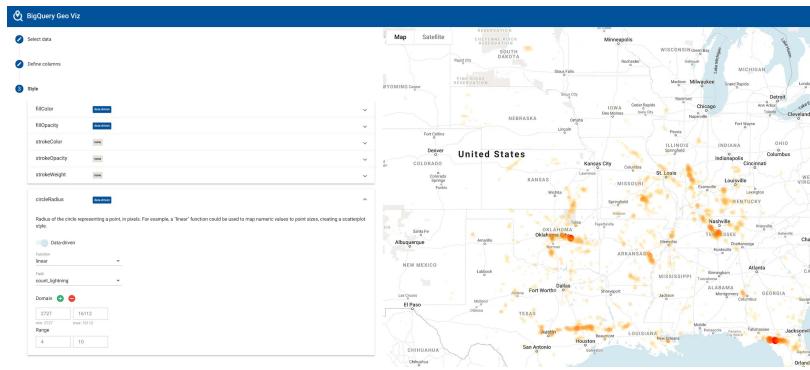
If your data is stored in JSON, use `ST_GeogFromGeoJSON`

```
SELECT
z.zip_code,
COUNT(*) AS num_stations
FROM
`bigquery-public-data.new_york_citibike.citibike_stations` AS s,
`bigquery-public-data.geo_us_boundaries.zip_codes` AS z
WHERE
ST_DWithin(z.zcta_geom,
ST_GeogPoint(s.longitude, s.latitude),
1000) -- 1km
AND num_bikes_available > 30
GROUP BY
z.zip_code
ORDER BY
num_stations DESC
```

```
SELECT ST_GeogFromGeoJSON('{"type": "Point", "coordinates": [-73.967416,40.756014]}'")
```



Visualization with BigQuery Geo Viz



Demo

Mapping Fastest Growing Zip
Codes with BigQuery GeoViz

Demo:

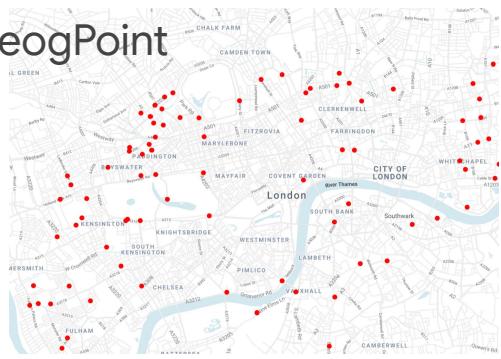
<https://github.com/GoogleCloudPlatform/training-data-analyst/tree/master/courses/data-engineering/demos>

Represent points with ST_GeogPoint

- Represented as WKT (Well Known Text)

```
SELECT
  id,
  longitude,
  latitude,
  ST_GeogPoint(longitude, latitude) AS location
FROM
  `bigquery-public-data.london_bicycles.cycle_stations`
LIMIT
  100
```

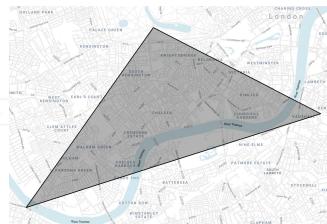
		longitude	latitude	location
471		-0.186753859	51.4916156	POINT(-0.186753859 51.4916156)
165		-0.183716959	51.517932	POINT(-0.183716959 51.517932)
261		-0.191351186	51.5134991	POINT(-0.191351186 51.5134991)
131		-0.136792671	51.53300545	POINT(-0.136792671 51.53300545)
467		-0.030556	51.523538	POINT(-0.030556 51.523538)
43		-0.15718945	51.52026	POINT(-0.15718945 51.52026)
212		-0.199004026	51.50958458	POINT(-0.199004026 51.50958458)
517		-0.033085	51.532513	POINT(-0.033085 51.532513)
704		-0.202802098	51.45682071	POINT(-0.202802098 51.45682071)
721		-0.026362677	51.53603947	POINT(-0.026362677 51.53603947)



Represent regions with ST_MakeLine and ST_MakePolygon

```
WITH stations AS (
SELECT
  (SELECT ST_GeogPoint(longitude, latitude) FROM `bigquery-public-data.london_bike_stations` WHERE location_id = 300) AS loc300,
  (SELECT ST_GeogPoint(longitude, latitude) FROM `bigquery-public-data.london_bike_stations` WHERE location_id = 305) AS loc305,
  (SELECT ST_GeogPoint(longitude, latitude) FROM `bigquery-public-data.london_bike_stations` WHERE location_id = 302) AS loc302
)
SELECT
  ST_MakeLine(loc300, loc305) AS seg1,
  ST_MakePolygon(ST_MakeLine([loc300, loc305, loc302])) AS poly
FROM
  stations
```

```
seg1
poly
LINESTRING(0.17306032 51.505014, -0.115853961 51.48677988)
POLYGON((0.216573 51.466907, -0.115853961 51.48677988, -0.17306032 ...))
```

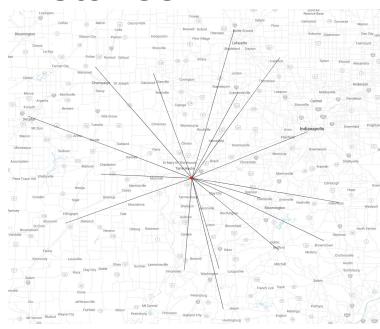


Are Locations Within Some Distance?

- ST_DWithin

```
SELECT
    m.name AS city,
    m.int_point AS city_coords,
    ST_MakeLine(
        n.int_point,
        m.int_point
    ) AS segs
FROM
    `bigquery-public-data.geo_us_boundaries.us_msas` AS n,
    `bigquery-public-data.geo_us_boundaries.us_msas` AS m
WHERE
    n.name='Terre Haute, IN'
    AND ST_DWithin(
        n.int_point,
        m.int_point,
        1.568) --150km
```

city	city_coords	segs
Crawfordsville, IN	POINT(-86.8927145 40.0402962)	LINESTRING(-87.3470958 39.392389, -86.892714...
Decatur, IL	POINT(-88.9615288 39.8602372)	LINESTRING(-87.3470958 39.392389, -88.961528...
Washington, IN	POINT(-87.076944 38.694689)	LINESTRING(-87.3470958 39.392389, -87.076944...
Indianapolis-Carmel-Anderson, IN	POINT(-86.2045408 39.7449323)	LINESTRING(-87.3470958 39.392389, -86.204540...
Lafayette-West Lafayette, IN	POINT(-86.9304747 40.5147171)	LINESTRING(-87.3470958 39.392389, -86.930474...
Bloomington, IN	POINT(-86.6717544 39.2417362)	LINESTRING(-87.3470958 39.392389, -86.671754...
Seymour, IN	POINT(-86.0425161 38.9119571)	LINESTRING(-87.3470958 39.392389, -86.042516...
Frankfort, IN	POINT(-86.4775665 40.305944)	LINESTRING(-87.3470958 39.392389, -86.477566...
Charleston-Mattoon, IL	POINT(-88.2422121 39.4231628)	LINESTRING(-87.3470958 39.392389, -88.242212...
Jasper, IN	POINT(-87.0359535 38.3841599)	LINESTRING(-87.3470958 39.392389, -87.035953...



Other predicate functions

- Do locations intersect?
 - ST_Intersects
- Is one geometry contained inside another?
 - ST_Contains
- Does a geography engulf another?
 - ST_CoveredBy

```
WITH geos AS (
  SELECT
    (SELECT state_geom FROM `bigquery-public-data.geo_us_boundaries.us_states`  

     WHERE state_name='Massachusetts') AS ma_poly,  

    (SELECT msa_geom FROM `bigquery-public-data.geo_us_boundaries.us_msa`  

     WHERE name='Boston-Cambridge-Newton, MA-NH') AS boston_poly,  

    (SELECT msa_geom FROM `bigquery-public-data.geo_us_boundaries.us_msa`  

     WHERE name='Seattle-Tacoma-Bellevue, WA') AS seattle_poly
)
SELECT
  ST_Intersects(boston_poly, ma_poly) boston_in_ma,
  ST_Intersects(seattle_poly, ma_poly) seattle_in_ma
FROM
  geos
```

boston_in_ma	seattle_in_ma
true	false



Distance between points

- ST_Distance
- Can lessen resolution of geo coordinates with ST_SnapToGrid
 - ST_SnapToGrid(pt, 0.01) rounds off the latitude and longitude of pt to the second decimal place.

```
WITH geos AS (
  SELECT
    (SELECT int_point FROM `bigquery-public-data.geo_us_boundaries.us_msa`
     WHERE name='Ann Arbor, MI') AS ann_arbor_loc,
    (SELECT int_point FROM `bigquery-public-data.geo_us_boundaries.us_msa`
     WHERE name='Columbus, OH') AS columbus_loc
)
SELECT
  [ST_Distance](ann_arbor_loc, columbus_loc)/1000 AS dist_in_km
FROM
  geos
  
```

dist_in_km

267.63049397704333

```
SELECT
  name,
  int_point AS point_full,
  [ST_SnapToGrid](int_point, .01) AS point_2decs
FROM
  `bigquery-public-data.geo_us_boundaries.us_msa`
WHERE
  name='Ann Arbor, MI'
```

name	point_full	point_2decs
Ann Arbor, MI	POINT(-83.8446336 42.2523272)	POINT(-83.84 42.25)



Combine multiple polygons into a single unit

- ST_Union_Agg
- Find centroid of the union with
ST_Centroid_Agg

```
SELECT
  zip_code,
  zcta_geom AS loc
FROM
  `bigquery-public-data.geo_us_boundaries.us_zip_codes`
WHERE
  city = 'Seattle city'
  AND state_code = 'WA'

SELECT
  [ST_UNION_ALL(zcta_geom)] as all_zips,
  [ST_CENTROID_ALL(zcta_geom)] as centroid
FROM
  `bigquery-public-data.geo_us_boundaries.us_zip_codes`
WHERE
  city = 'Seattle city'
  AND state_code = 'WA'
```



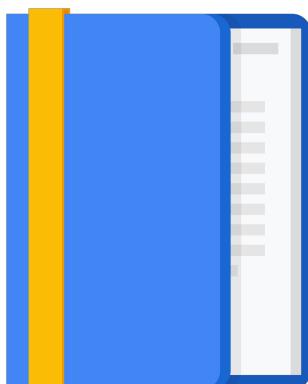
Agenda

Analytic Window Functions

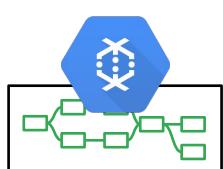
Using With Clauses

GIS Functions

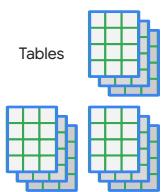
Performance Considerations



Best practices for fast, smart, data-driven decisions



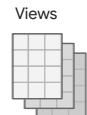
Use Dataflow to do the processing and data transformations



Create multiple tables for easy analysis



Use BigQuery for streaming analysis and dashboards
Store data in BigQuery for low cost long term storage



Create views for common queries



This is the goal: fast smart decisions.

Best practices for analyzing data with BigQuery

Dataset



Know your data

Questions



Ask good questions

SQL

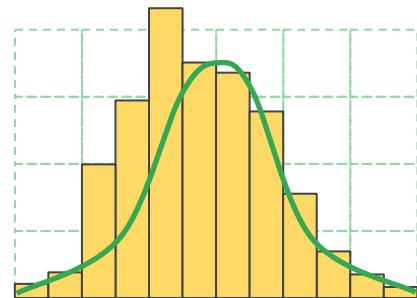
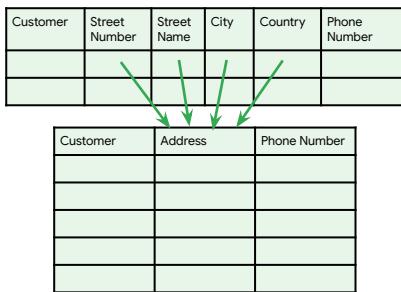
```
SELECT  
    name,  
    revenue  
FROM dataset  
ORDER BY revenue;
```

Use queries effectively



Exploring a dataset through SQL is more than just writing good code. You need to know what destination you're heading towards and the general layout of your data. Good data analysts will explore how the dataset is structured even before writing a single line of code.

How to optimize in production? Revisit the schema. Revisit the data.



Stop accumulating work that could be done earlier.



People often analyze data and develop a schema at the beginning of a project and never revisit those decisions. The assumptions they made at the beginning may have changed and are no longer true. So they attempt to adjust the downstream processes without ever reviewing and considering changing some of the original decisions.

Look at the data. Perhaps it was evenly distributed at the start of the project but as the work has grown, the data may have become skewed.

Look at the schemas. What were the goals then? Are those the same goals now? Is the organization of the data optimized for current operations?

Stop accumulating work that could be done earlier. Analogy: dirty dishes. If you clean them as you use them, the kitchen remains clean. If you save them, you end up with a sink full of dirty dishes and a lot of work.

<INSTRUCTOR>

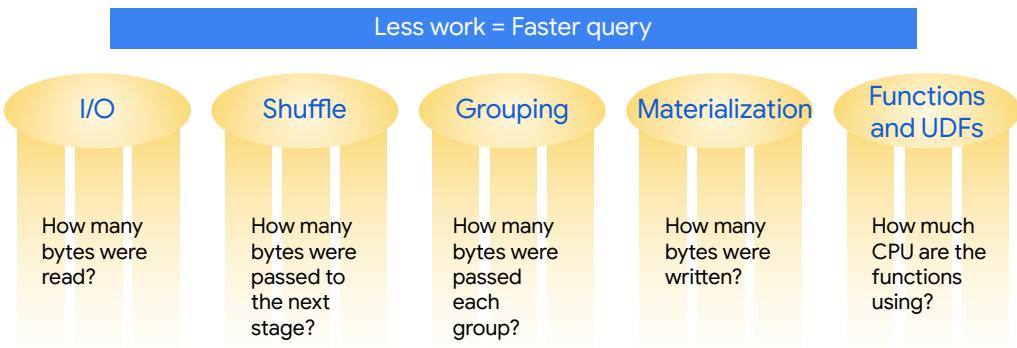
True story. I was working on an application for an organization. At first the data was exploratory. We went through a complicated schema design process that captured many of the nuances of the circumstance in a complex collection of relational tables. After weeks of work, the Data Analysts built a report generation system that was amazingly complex and took four days to run. Unfortunately, in production, the report was needed every two days. A consultant company was hired. After weeks of analysis they suggested replicating the entire system so that reports could be run in parallel to meet the schedule requirements.

I started experimenting with optimizing the system. I started with intermediate materialization of tables. And that made the report generation faster. Then I realized that we could get the report out much faster if we stored the RESULTS the report needed. I created intermediate tables with the fields that were needed for the last stages of the report. Eventually I rewrote the data ingress portion of the application. Instead of accumulating data and processing it, I pushed the processing back to data ingest.

For example, we had a total number of items in the final report. And in the original application, this required counting several million items. I created a table with the final count of items in it. And wrapped the operator functions so that every time an object was added to the database this field in the intermediate table was incremented, and each time an object was removed this field was decremented. Now the report already had this sum total when it began its queries and didn't need to calculate it dynamically. The time to run the report was radically reduced without buying a second system.

</INSTRUCTOR>

Improve scalability by improving efficiency



Old Silicon Valley saying: "Don't scale up your problems. Solve them early while they are small."

Optimize BigQuery queries

SELECT *

Avoid using unnecessary columns

WHERE

Filter early and often

JOIN

Put the largest table on the left

GROUP BY

Low cardinality is faster than high cardinality

APPROX_COUNT_DISTINCT

Is faster than

COUNT(DISTINCT)

ORDER

On the outermost query

*

Use wildcards to query multiple tables

--time_partitioning_type

Time partitioning tables for easier search



See: <https://cloud.google.com/bigquery/docs/best-practices-performance-compute>

BigQuery supports three ways of partitioning tables

Ingestion time

```
bq query --destination_table mydataset.mytable  
--time_partitioning_type=DAY  
...
```

Any column that is of type
DATE or TIMESTAMP

```
bq mk --table --schema a:STRING,tm:TIMESTAMP  
--time_partitioning_field tm
```

Integer-typed column

```
bq mk --table --schema "customer_id:integer,value:integer"  
--range_partitioning=customer_id,0,100,10 my_dataset.my_table
```



If you remember, earlier in this course we talked about how to build a data warehouse. We mentioned that you can optimize the tables in your data warehouse by reducing the cost and amount of data read. This can be achieved by partitioning your tables.

Partitioning tables is very relevant to performance so let's revisit in the next few slides some of the main points already covered.

You enable partitioning during the table-creation process.

The slide shows how to migrate an existing table to an ingestion-time-partitioned table: just use destination table. It will cost you one table scan.

BigQuery creates new date-based partitions automatically, with no need for additional maintenance. In addition, you can specify an expiration time for data in the partitions.

New data that is inserted into a partitioned table is written to the raw partition at the time of insert. To explicitly control which partition the data is loaded to, your load job can specify a particular date partition.

Partitioning can improve query cost and performance by cutting down on data being queried

```
SELECT  
    field1  
FROM  
    mydataset.table1  
WHERE  
    _PARTITIONTIME > TIMESTAMP_SUB(TIMESTAMP('2016-04-15'), INTERVAL 5 DAY)
```

Isolate the partition field in the left-hand side of the query expression!

```
bq query \  
--destination_table mydataset.mytable  
--time_partitioning_type=DAY --require_partition_filter  
...
```



In a table partitioned by a date or timestamp column, each partition contains a single day of data. When the data is stored, BigQuery ensures that all the data in a block belongs to a single partition. A partitioned table maintains these properties across all operations that modify it: query jobs, Data Manipulation Language (DML) statements, Data Definition Language (DDL) statements, load jobs, and copy jobs. This requires BigQuery to maintain more metadata than a non-partitioned table. As the number of partitions increases, the amount of metadata overhead increases.

Although more metadata must be maintained, by ensuring that data is partitioned globally, BigQuery can more accurately estimate the bytes processed by a query before you run it. This cost calculation provides an upper bound on the final cost of the query.

A good practice is to require that queries always include the partition filter.

Make sure that the partition field is isolated on the left-hand-side since that's the only way BigQuery can quickly discard unnecessary partitions.

BigQuery automatically sorts the data based on values in the clustering columns

c1	c2	c3	eventDate	c5
Blue	Blue	Blue	2019-01-01	Blue
Blue	Blue	Blue	2019-01-02	Blue
Yellow	Yellow	Yellow	2019-01-03	Yellow
Yellow	Yellow	Yellow	2019-01-04	Yellow
Blue	Blue	Blue	2019-01-05	Blue

```
SELECT c1, c3 FROM ...
WHERE eventDate BETWEEN "2019-01-03" AND
"2019-01-04"
```

Partitioned tables



c1	userId	c3	eventDate	c5
Red	Red	Red	2019-01-01	Red
Red	Red	Red	2019-01-02	Red
Yellow	Red	Red	2019-01-03	Yellow
Yellow	Red	Red	2019-01-04	Yellow
Red	Red	Red	2019-01-05	Red

```
SELECT c1, c3 FROM ... WHERE userId BETWEEN 52
and 63 AND eventDate BETWEEN "2019-01-03" AND
"2019-01-04"
```

Clustered tables

Clustering can improve the performance of certain types of queries, such as queries that use filter clauses and those that aggregate data. When data is written to a clustered table by a query or load job, BigQuery sorts the data using the values in the clustering columns. These values are used to organize the data into multiple blocks in BigQuery storage. When you submit a query containing a clause that filters data based on the clustering columns, BigQuery uses the sorted blocks to eliminate scans of unnecessary data.

Similarly, when you submit a query that aggregates data based on the values in the clustering columns, performance is improved because the sorted blocks co-locate rows with similar values.

In this example, the table is partitioned by eventDate and clustered by userId. Now, because the query looks for partitions in a specific range, only 2 of the 5 partitions are considered.

Because the query looks for userId in a specific range, BigQuery can jump to the row range and read only those rows for each of the columns needed.

Set up clustering at table creation time

c1	userId	c3	eventDate	c5
Blue	Red	Blue	2019-01-01	Blue
Blue	Red	Blue	2019-01-02	Blue
Yellow	Red	Yellow	2019-01-03	Yellow
Yellow	Red	Yellow	2019-01-04	Yellow
Blue	Red	Blue	2019-01-05	Blue

```
CREATE TABLE mydataset.myclusteredtable
(
  c1 NUMERIC,
  userId STRING,
  c3 STRING,
  eventDate TIMESTAMP,
  c5 GEOGRAPHY
)
PARTITION BY DATE(eventDate)
CLUSTER BY userId
OPTIONS (
  partition_expiration_days=3,
  description="cluster")
AS SELECT * FROM mydataset.myothertable
```



You set up clustering at table creation time. Here, we are creating the table, partitioning by eventDate, and clustering by userId. We are also telling BigQuery to expire partitions that are more than 3 days old.

The columns you specify in the cluster are used to co-locate related data. When you cluster a table using multiple columns, the order of columns you specify is important. The order of the specified columns determines the sort order of the data.

In streaming tables, the sorting fails over time, and so BigQuery has to recluster

```
UPDATE ds.table  
SET c1 = 300  
WHERE c1 = 300  
AND eventDate > TIMESTAMP_SUB(CURRENT_TIMESTAMP(), INTERVAL 1 DAY)
```

Can force a recluster using DML on necessary partition

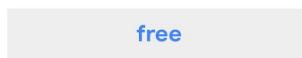
c1	userId	c3	eventDate	c5
blue	red	blue	2019-01-01	blue
blue	red	blue	2019-01-02	blue
yellow	red	yellow	2019-01-03	yellow
yellow	red	yellow	2019-01-04	yellow
blue	red	blue	2019-01-05	blue



Over time, as more and more operations modify a table, the degree to which the data is sorted begins to weaken, and the table becomes only partially sorted. In a partially sorted table, queries that use the clustering columns may need to scan more blocks compared to a table that is fully sorted. You can re-cluster the data in the entire table by running a `SELECT *` query that selects from and overwrites the table, but guess what! You don't need to do that anymore.

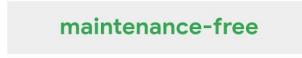
BigQuery will automatically recluster your data

Automatic re-clustering



free

Doesn't consume your query resources



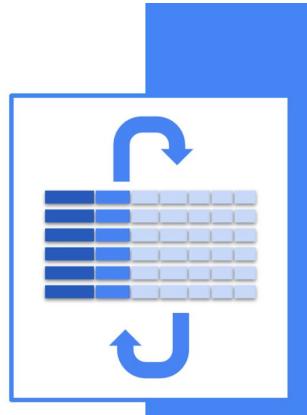
maintenance-free

Requires no setup or maintenance



autonomous

Automatically happens in the background



The great news is that BigQuery now periodically does auto-reclustering for you so you don't need to worry about your clusters getting out of date as you get new data.

Automatic re-clustering is absolutely free and automatically happens in the background -- you don't need to do anything additional to enable it.

<https://cloud.google.com/blog/products/data-analytics/whats-happening-bigquery-adding-speed-and-flexibility-10x-streaming-quota-cloud-sql-federation-and-more>
https://cloud.google.com/bigquery/docs/clustered-tables#automatic_re-clustering

Organize data through managed tables

Partitioning

Filtering storage before query execution begins to reduce costs.

Reduces a full table scan to the partitions specified.

A single column results in lower cardinality (e.g., thousands of partitions).

- Time partitioning (Pseudocolumn)
- Time partitioning (User Date/Time column)
- Integer range partitioning

Clustering

Storage optimization within columnar segments to improve filtering and record colocation.

Clustering performance and cost savings can't be assessed before query begins.

Prioritized clustering of up to 4 columns, on more diverse types (but no nested columns).



Partitioning provides a way to obtain accurate cost estimates for queries and guarantees improved cost and performance. Clustering provides additional cost and performance benefits in addition to the partitioning benefits.

When to use clustering



Your data is already partitioned on a DATE or TIMESTAMP or Integer Range



You commonly use filters or aggregation against particular columns in your queries.

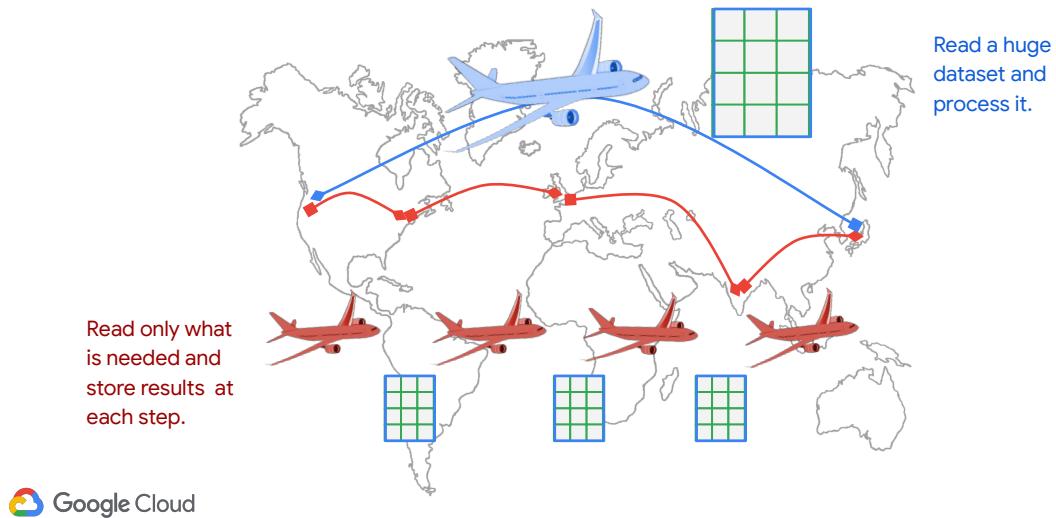


BigQuery supports clustering for both partitioned and non-partitioned tables.

When you use clustering and partitioning together, the data can be partitioned by a date or timestamp column and then clustered on a different set of columns. In this case, data in each partition is clustered based on the values of the clustering columns. Partitioning provides a way to obtain accurate cost estimates for queries.

Keep in mind if you don't have partitioned columns and you want the benefits of clustering you can create a `fake_date` column of type DATE and have all the values be NULL.

Break queries into stages using intermediate tables



If you create a large, multi-stage query, each time you run it, BigQuery reads all the data that is required by the query. All the data that is read each time the query is run.

Intermediate table materialization is where you break the query into stages. Each stage materializes the query results by writing them to a destination table.

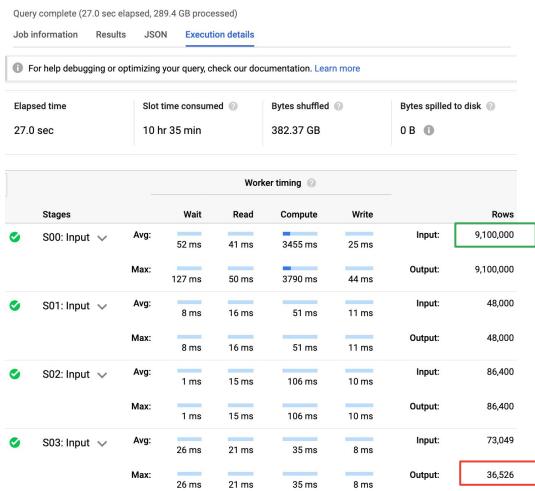
Querying the smaller destination table reduces the amount of data that is read. In general, storing the smaller materialized results is more efficient than processing the larger amount of data.

The analogy is air travel from Sunnyvale, California, USA to Japan. There is one direct flight. Or a series of four shorter connecting flights. The direct flight has to carry the fuel for the entire journey. The connecting flights only need enough fuel for each leg of the trip. The total fuel used in landing and taking off (an analogy for storing the intermediate tables) was less than the total fuel used for carrying everything in the entire journey.

Here is a tip: Compare costs of storing the data with costs of processing the data. Processing the large dataset will use more processing resources. Storing the intermediate tables will use more storage resources. In general, processing data is more expensive than storing data. But you can do the calculations yourself to establish a breakeven for your particular use case.

<https://pixabay.com/vectors/airplane-jet-aircraft-plane-flight-309503/>

Track input and output counts with Execution details tab

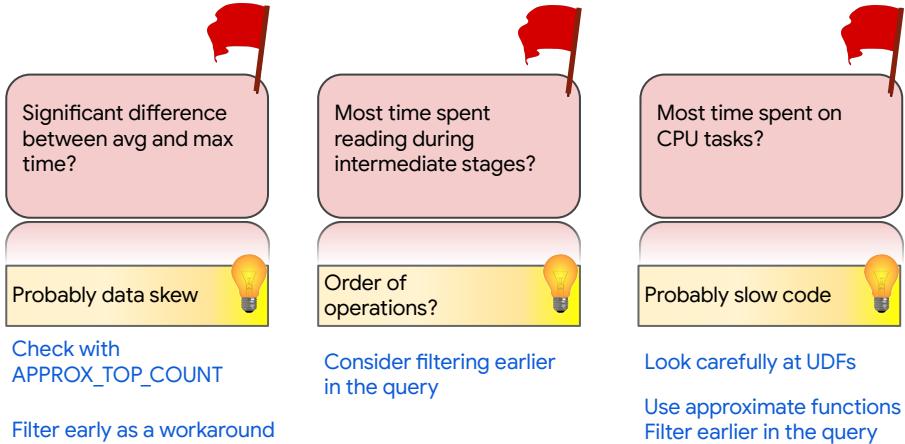


A different way to check how many records are being processed is by clicking on the Explanation tab in the BigQuery UI after running a query.

We started with 9.1 Million rows and filtered down to roughly 36k

The query stages represent how BigQuery mapped out the work required to perform the query job.

Using BigQuery plans to optimize



Approximate Functions: APPROX_COUNT_DISTINCT(expression) Description.
Returns the approximate result for COUNT(DISTINCT expression). The result is less accurate but it performs much more efficiently.

Analyze BigQuery performance in Cloud Monitoring

Custom Dashboards

Metrics include:

- slots utilization
- queries in flight
- upload bytes
- stored bytes



These charts show Slot Utilization, Slots available and queries in flight for a 1 hr period.



Optimizing your BigQuery Queries for Performance

Objectives

- Use BigQuery to
 - Minimize I/O
 - Cache results of previous queries
 - Perform efficient joins
 - Avoid overwhelming single workers
 - Use approximate aggregation functions

Optimizing your BigQuery Queries for Performance

<https://gcpstaging.qwiklabs.com/labs/24799/edit>

Once your data is loaded into BigQuery, you are charged for storing it

Active Storage Pricing

Storage pricing is prorated per MB, per second. For example, if you store:

- 100 MB for half a month, you pay \$0.001 (a tenth of a cent)
- 500 GB for half a month, you pay \$5 (\$0.02/GB per month)
- 1 TB for a full month, you pay \$20

Long-term Storage Pricing

- Table or partition not edited 90+ consecutive days
- Pricing drops \approx 50%
- No degradation (performance, durability, availability, functionality)
- Applies to BigQuery storage only



Storage pricing is based on the amount of data stored in your tables when it is uncompressed. The size of the data is calculated based on the data types of the individual columns. Active storage pricing is prorated per MB, per second.

If a table is not edited for 90 consecutive days it is considered long-term storage. The price of storage for that table automatically drops by approximately 50 percent. There is no degradation of performance, durability, availability, or any other functionality.

If the table is edited, the price reverts back to the regular storage pricing, and the 90-day timer starts counting from zero. Anything that modifies the data in a table resets the timer, including:

- Loading data into a table
- Copying data into a table
- Writing query results to a table
- Using the Data Manipulation Language
- Using Data Definition Language
- Streaming data into the table

All other actions do not reset the timer, including:

- Querying a table
- Creating a view that queries a table
- Exporting data from a table
- Copying a table (to another destination table)
- Patching or updating a table resource

<https://cloud.google.com/bigquery/pricing>

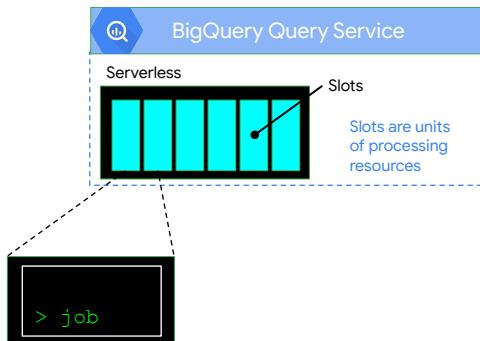
Which BigQuery pricing model to pick?

On-Demand Pricing	Flat-Rate Pricing
<ul style="list-style-type: none">• \$5/TB of data processed• Quota limit of 2,000 slots• Slots shared amongst all on demand users• 1st TB of data processed is free each month	<ul style="list-style-type: none">• Fixed rate pricing is \$10k* per 500 slots per month• Slots are dedicated 24/7• Starting at 500 slots• Unlimited use of slots
Good for exploratory work and discovery	Good for multiple large workloads and consistent monthly billing



Fixed rate pricing is \$10k* per 500 slots per month. A 25% discount is offered for customers choosing a term length of at least 1-year (\$7.5k per 500 slots). Capacity is sold in increments of 500 slots with a current minimum of 500 slots.

Slots are units of resources that are consumed when a query is run



Because you don't get to see the VMs, BigQuery exposes slots to help you manage resource consumption and costs. Slots are units of resources such as processing power and memory that are consumed when a query is run. BigQuery automatically calculates how many slots are required by each query, depending on query size and complexity. The default slot capacity and allocation work well in most cases. You can monitor slot usage in Stackdriver.

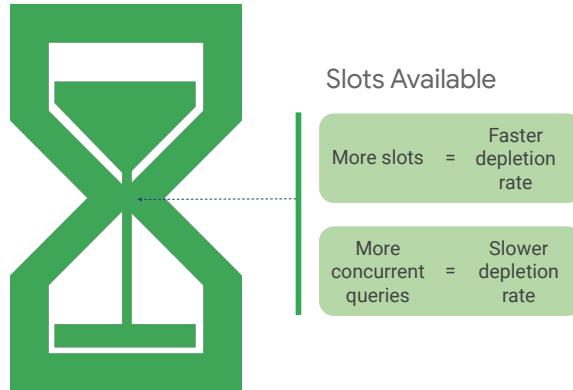
Candidate circumstances where additional slot capacity might improve performance are solutions with very complex queries on very large datasets with highly concurrent workloads. You can read more about slots in the online documentation or contact a sales representative.

Estimating the right BigQuery slots allocation is critical

Guideline: It is recommended to plan for 2,000 BigQuery slots for every 50 medium-complexity queries simultaneously

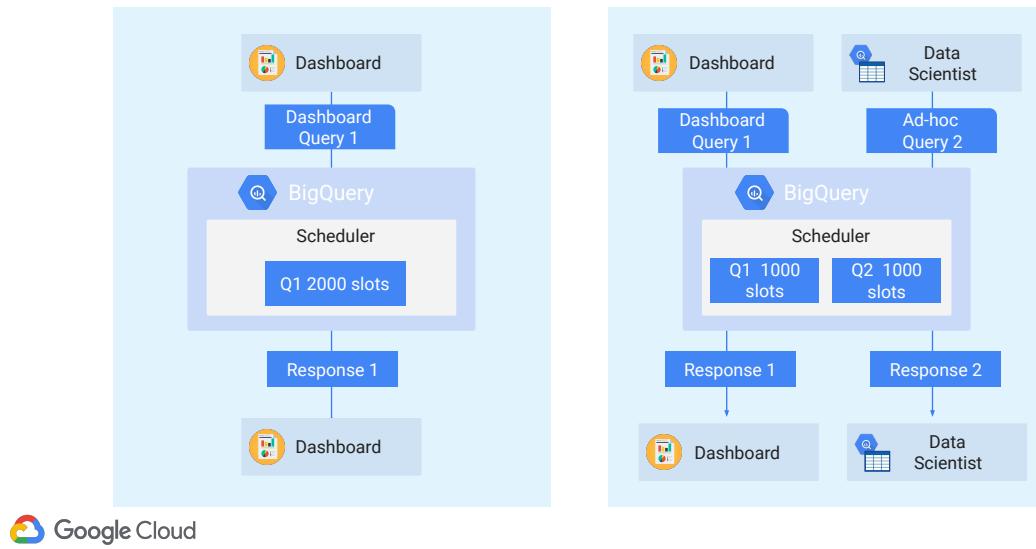
50 queries = 2,000 Slots

(quota: 100 concurrent queries, but this can be raised)



BigQuery doesn't support fine-grained prioritization of interactive or batch queries. To avoid a pile up of BigQuery jobs and timely execution, estimating the right BigQuery slots allocation is critical. Currently, BQ times out any query taking longer than 6 hours.

BigQuery has a fair scheduler



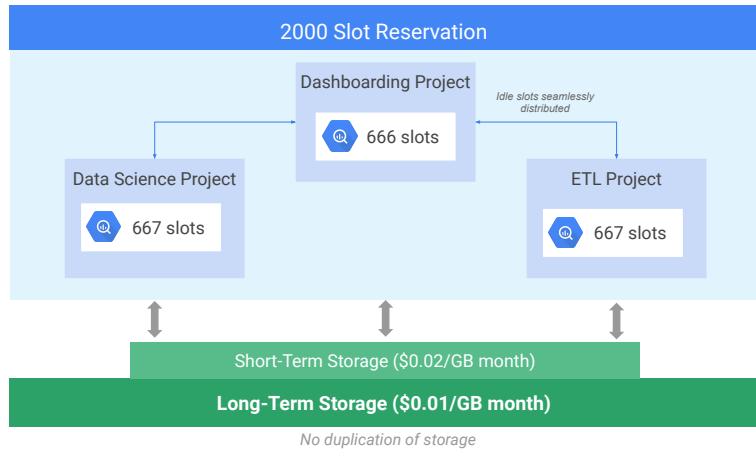
If one query is executing within BigQuery, it has full access to the amount of slots available to the project or reservation, by default 2000

If we suddenly execute a second query, BigQuery will split the slots between the 2 queries with each getting half the total amount of slots available, in this case 1000 each

This subdividing of compute resources will continue to happen as more queries are executed

This is a long way of saying, it's unlikely that one resource-heavy query will overpower the system and steal resources from other running queries

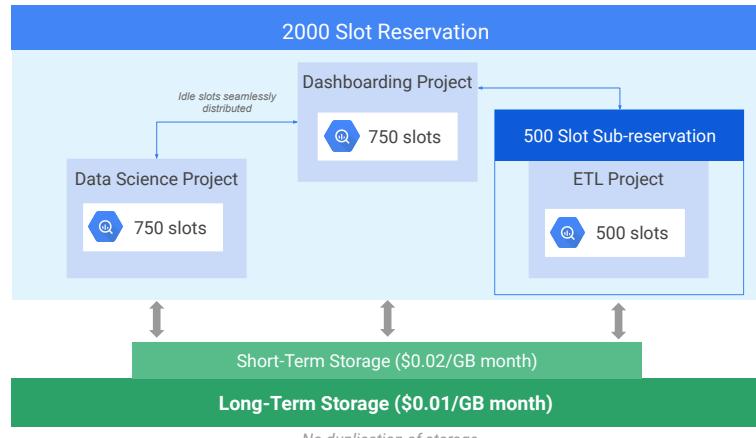
Concurrency is fair across projects, users, and queries, with unused capacity fairly divided among existing tasks



In flat rate pricing they will have a fixed number of slots. Concurrency is fair across projects, users, and queries. That is if you have 2k slots and 2 projects, each project can get up to 1k slots. If one project uses less, the other project will be able to use all of the remainder. If you have 2 users in each project, each user will be able to get 500 slots. And if each of the two users each of the two projects, runs two queries, they'll each get 500 slots.

This is a long way of saying they won't likely degrade performance by adding projects.

To prioritize projects, set up a hierarchical reservation



Note that if they want to prioritize one project over another, they can set up a hierarchical reservation. Let's say you have an ETL project that is somewhat lower priority than your dashboarding project. You can give the ETL project 500 slots as a sub-reservation and the dashboarding project will be in the outer one. If both projects are fully using their reservations, the ETL project can never get more than 500 slots. When one project is lightly used, the other project will be able to take the remaining slots.



Creating Date-Partitioned Tables in BigQuery

Objectives

- Query a partitioned dataset
- Create dataset partitions to improve query performance and reduce cost

OPTIONAL LAB

Creating Date-Partitioned Tables in BigQuery

Module Summary

- Use BigQuery's analytic window functions for advanced analysis
- BigQuery has built in GIS functionality
- Use BigQuery's WITH clause to help modularize your queries
- There are many ways to optimize your query performance

